

POSTGRESQL IS YESQL!

All Your Base Conference 2015

Dimitri Fontaine dimitri@2ndQuadrant.fr
@tapoueh

13 November 2015

Dimitri Fontaine

PRINCIPAL CONSULTANT AT 2NDQUADRANT

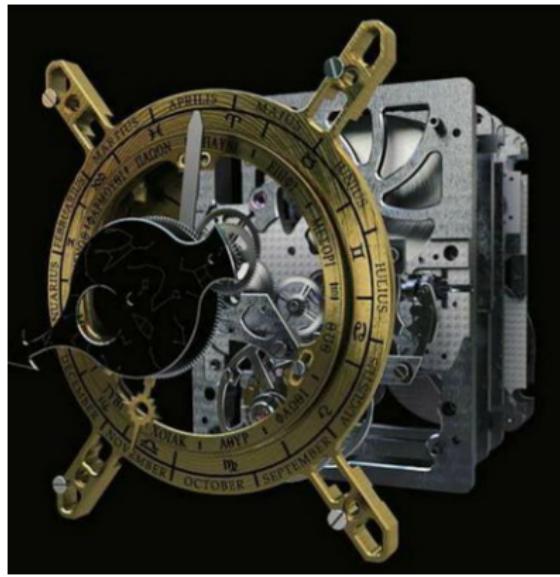


Relational Database Management System



Antikythera mechanism

Relational Database Management System

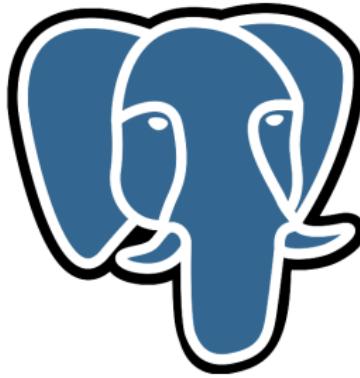


Antikythera mechanism

Relational Database Management System



POSTGRESQL IS YESQL!



POSTGRESQL IS YESQL!

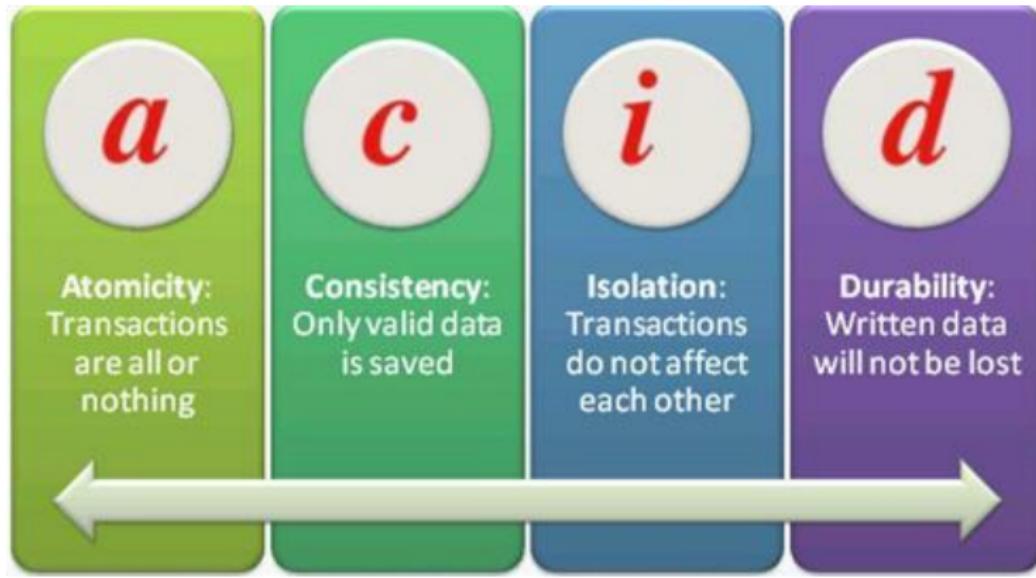


Not Only SQL

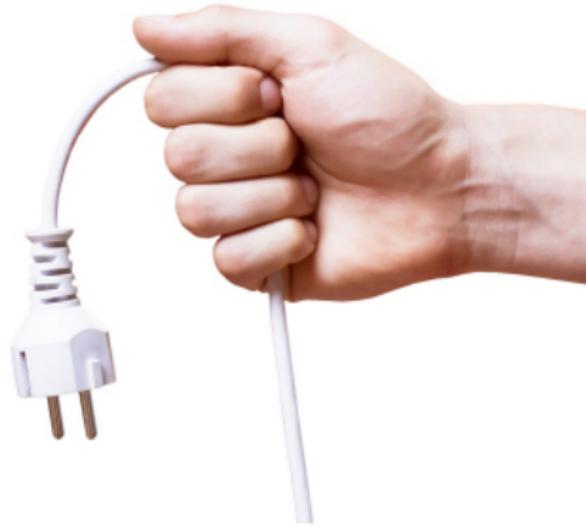
Relational Database Management System



Relational Database Management System



D is for Durability



A is for Atomic, that is,
Transactions

Rollback

Transactions

```
BEGIN;  
  DROP TABLE critical_data;  
ROLLBACK;
```

I is for Isolated, that is,
Transactions



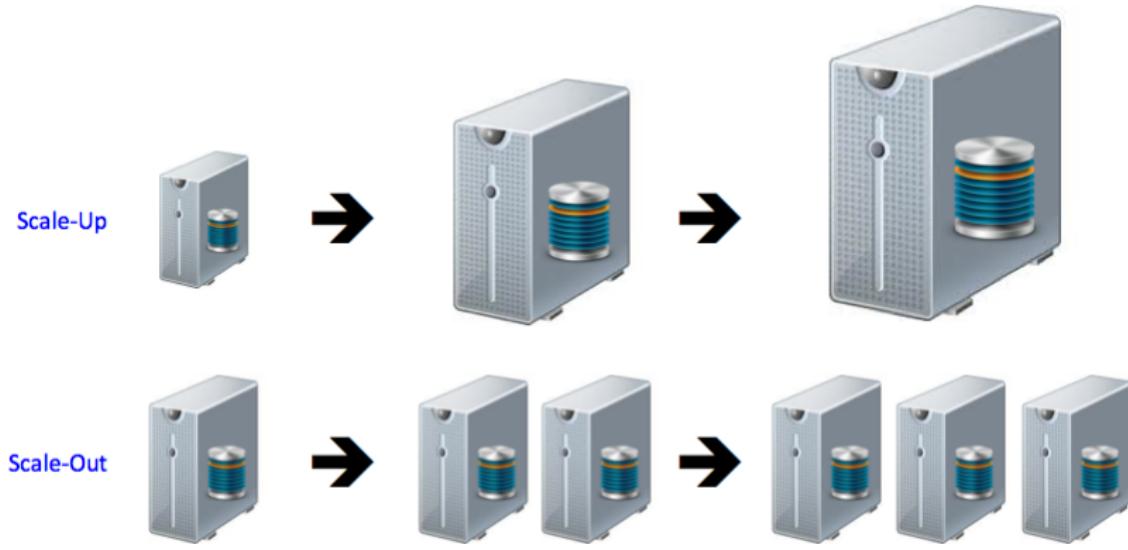
You don't want backups, you want data loss recovery!



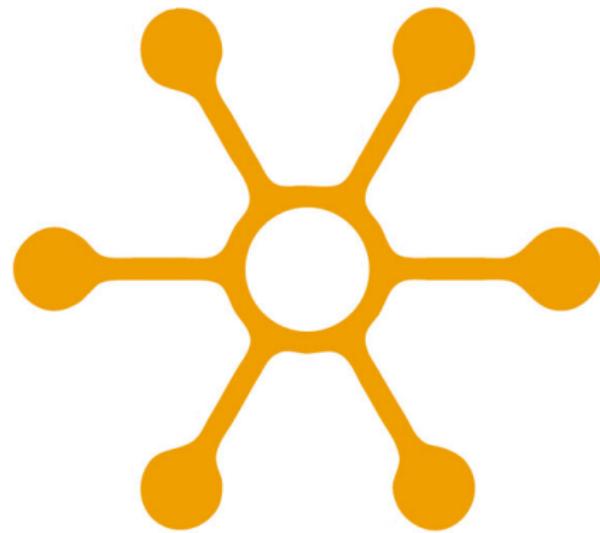
C is for Consistency



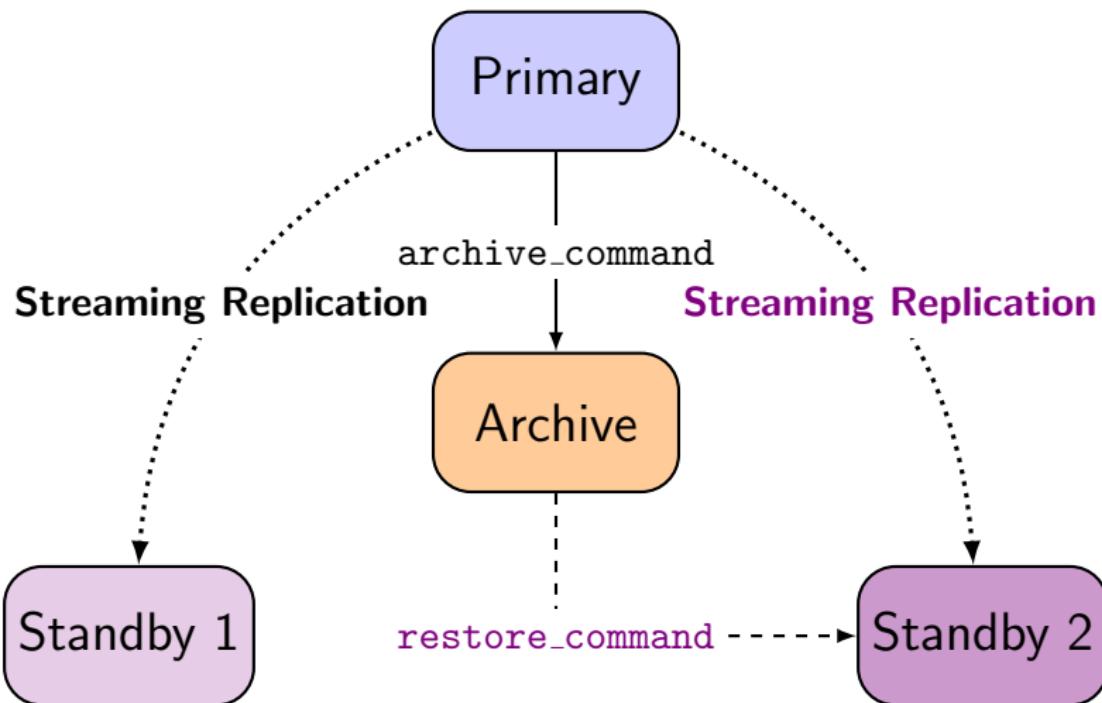
Scaling: up or out?



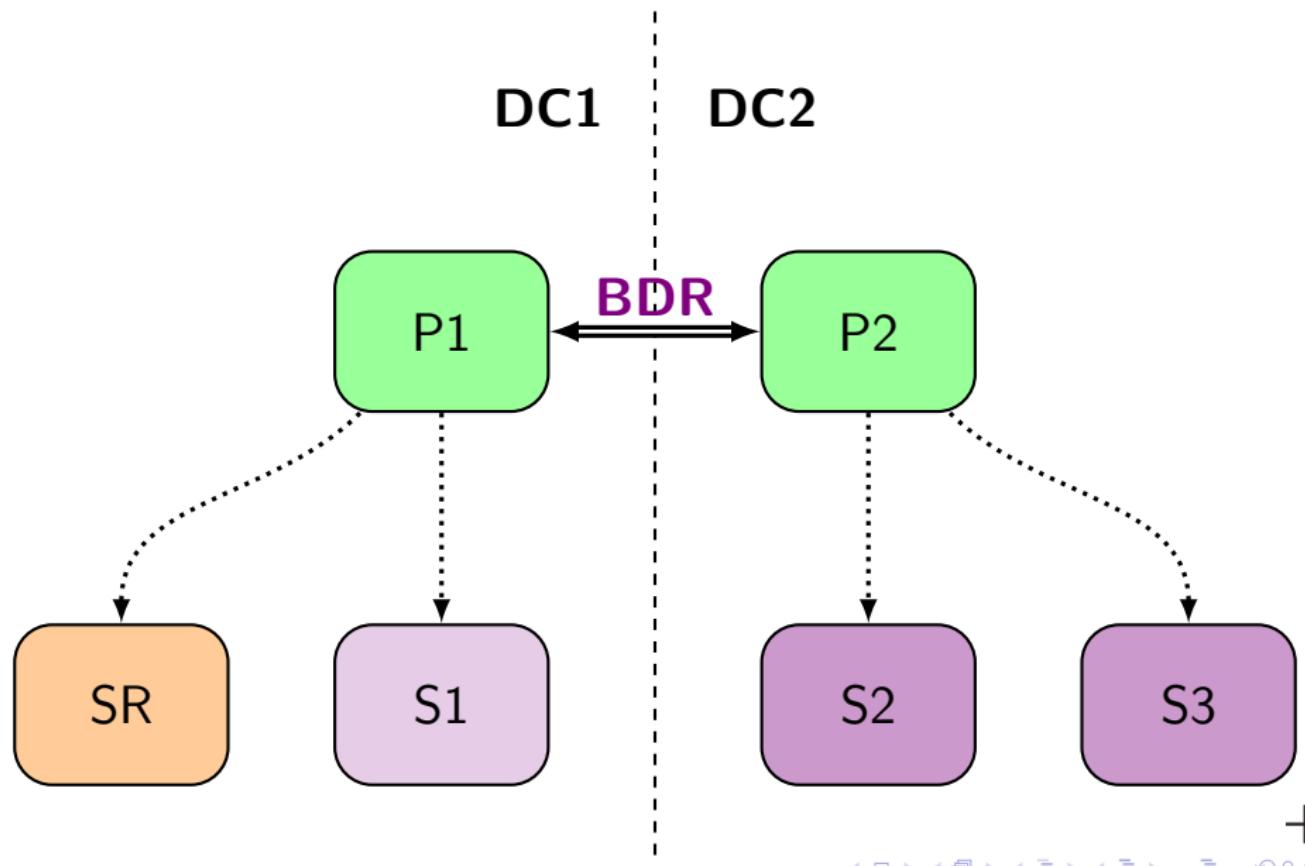
High Availability of Services and Data



High Availability: Consistency and Durability



High Availability



Consistency



Data Types and Constraints

Data integrity First Layer: data type

- Integer
- Arbitrary precision numbers, UUID
- Floating point
- Character, Text
- Bytea, bitstring
- Date/Time, Time Zones
- Boolean
- Enum, Arrays, Composite Types, Range Types
- Point, Line Segments, Boxes
- Paths, Polygons, Circles
- Inet, CIDR, Macaddr
- JSON, XML

Advanced data types and constraints

```
CREATE TABLE reservation
(
    room      text,
    professor text,
    during    period,
    EXCLUDE USING gist
    (
        room with =,
        during with &&
    )
);
```

```
CREATE TABLE circles (
    c circle,
    EXCLUDE USING gist
    (c WITH &&)
);
```

STRUCTURED QUERY LANGUAGE



Finding the last counter value before reset

Write some SQL here

tick	nb	max
1	0	
2	10	
3	20	
4	30	
5	40	40
6	0	
7	20	
8	30	
9	60	60

(9 rows)

Window Functions: lead() over()

```
select tick,
       nb,
       lead(nb) over (order by tick)
  from measures;
```

tick	nb	lead
1	0	10
2	10	20
3	20	30
4	30	40
5	40	0
6	0	20
7	20	30
8	30	60
9	60	

(9 rows)

Window Functions and CASE

```
select tick, nb,
       case when lead(nb) over w < nb
             then nb
             when lead(nb) over w is null
             then nb
             else null
         end as max
  from measures
 window w as (order by tick);
```

tick	nb	max
1	0	
2	10	
3	20	
4	30	
5	40	40
6	0	
7	20	
8	30	
9	60	60

(9 rows)

PostgreSQL Extensibility

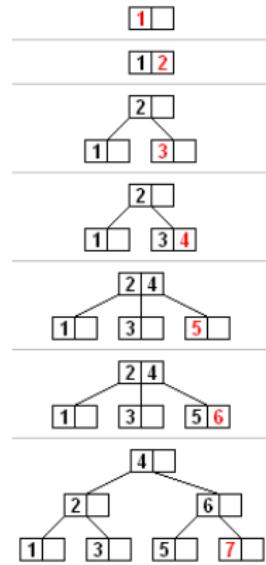
PostgreSQL is *highly extensible*



B-Tree

Balanced Tree, the default index type

- Built for speed
- *unique* concurrency tricks
- Balanced
- support function: cmp
- operators: <= < = > >=

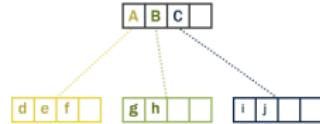
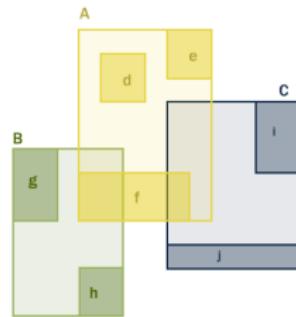


Generalized Index Search Tree

GiST or the Indexing API

- Built for comfort
- Balanced
- API: consistent, same, union
- API: penalty, picksplit
- API: compress, decompress
- operators: @> <@ && @@ = &< &>
<<| ...

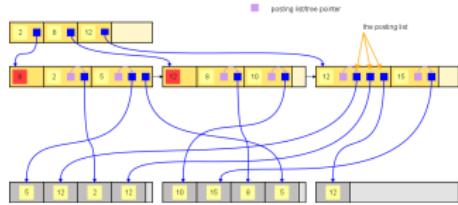
R-tree Hierarchy



Generalized Inverted iNdex

Indexing several pointers per value, inverted cardinality

- Built for Text Search and Arrays
- Balanced
- API: compare, consistent
- API: extractValue, extractQuery
- operators: @> <@ && =



PostgreSQL Full Text Search

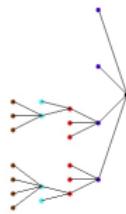
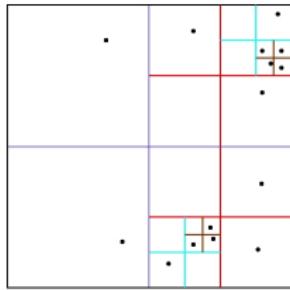


Spaced Partition GiST

Unbalanced index, used for **GIS**



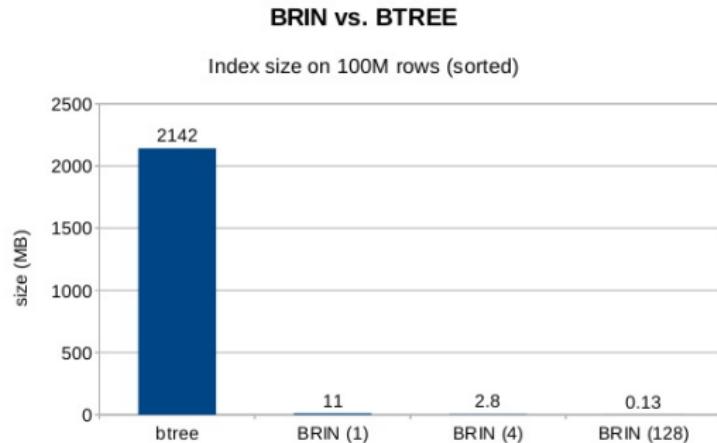
Quadtree



Oleg Bartunov, Teodor Sigaev

PGCon-2011, Ottawa, May 17-20, 2011

Block Range Index for Big Data



2ndQuadrant +
Professional PostgreSQL

STRUCTURED QUERY LANGUAGE



IP Ranges, ip4r, Geolocation

PostgreSQL allows using SQL and JOINs to match IP4R with geolocation.

```
select *
  from geolite.blocks
  join geolite.location
    using(locid)
  where iprange
        >>=
      '74.125.195.147' ;
```



IP Ranges, ip4r, Geolocation

PostgreSQL allows using SQL and JOINs to match IP4R with geolocation.

```
select *  
  from geolite.blocks  
  join geolite.location  
    using(locid)  
 where iprange  
       >=>  
     '74.125.195.147';
```

```
- [ RECORD 1 ] -----  
 locid      | 2703  
 iprange   | 74.125.189.24-74.125.  
 country    | US  
 region     | CA  
 city       | Mountain View  
 postalcode | 94043  
 location   | (-122.0574,37.4192)  
 metrocode   | 807  
 areacode    | 650
```

Time: 1.335 ms

Geolocation: ip4r meets earthdistance

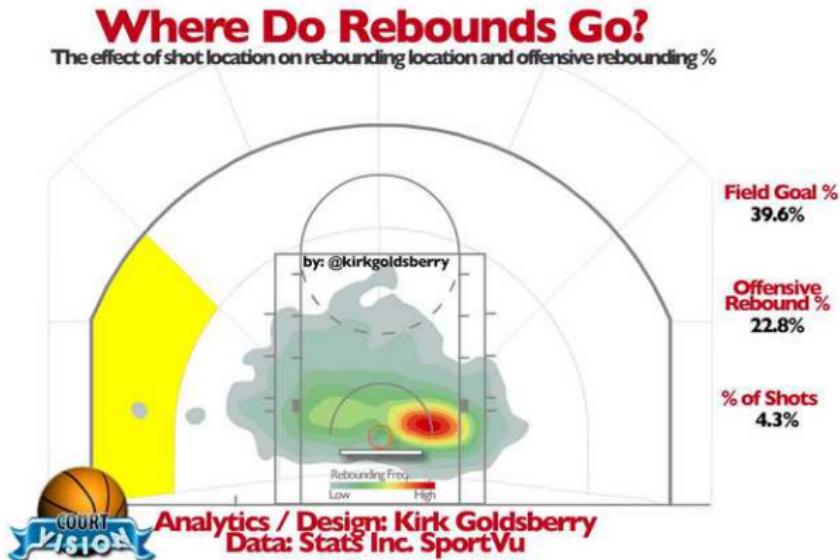


Some pubs nearby... some place...

```
with geoloc as
(
  select location as l
  from location
  join blocks using(locid)
  where iprange
    >>=
    '212.58.251.195'
)
select name,
       pos <@> 1 miles
  from pubnames, geoloc
 order by pos <-> 1
 limit 10;
```

name	miles
Blue Anchor	0.299
Dukes Head	0.360
Blue Ball	0.337
Bell (aka The Rat)	0.481
on the Green	0.602
Fox & Hounds	0.549
Chequers	0.712
Sportsman	1.377
Kingswood Arms	1.205
Tottenham Corner	2.007
(10 rows)	

Time: 3.275 ms



NBA Games Statistics

An interesting factoid: the team that recorded the fewest defensive rebounds in a win was the 1995-96 Toronto Raptors, who beat the Milwaukee Bucks 93-87 on 12/26/1995 despite recording only 14 defensive rebounds.

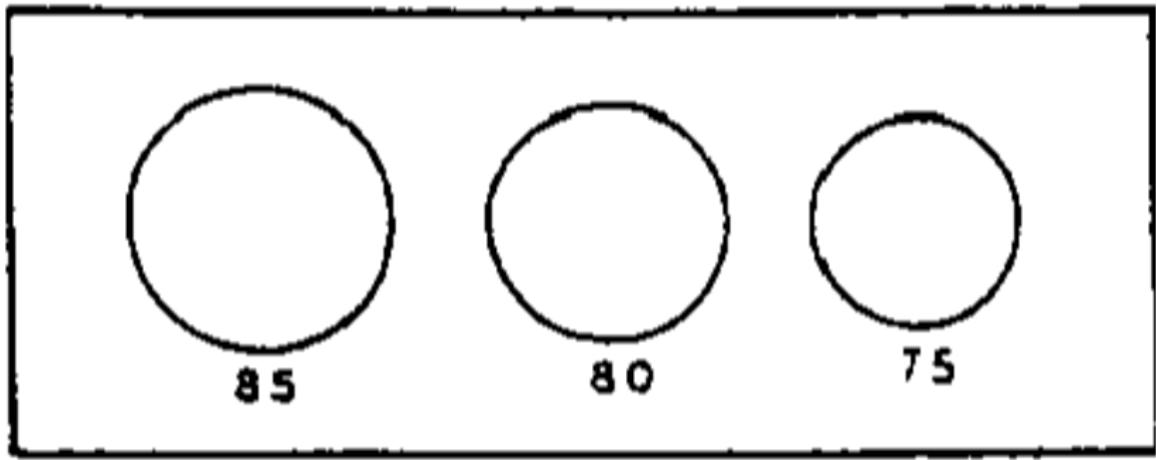
PostgreSQL Data Mining

```
with stats(game, team, drb, min) as (
    select ts.game, ts.team, drb, min(drb) over ()
        from team_stats ts
            join winners w on w.id = ts.game
                and w.winner = ts.team
)
select game.date::date,
    host.name || ' -- ' || host_score as host,
    guest.name || ' -- ' || guest_score as guest,
    stats.drb as winner_drb
from stats
    join game on game.id = stats.game
    join team host on host.id = game.host
    join team guest on guest.id = game.guest
where drb = min;
```

PostgreSQL Data Mining

```
-[ RECORD 1 ]-----  
date      | 1995-12-26  
host      | Toronto Raptors -- 93  
guest     | Milwaukee Bucks -- 87  
winner_drb | 14  
-[ RECORD 2 ]-----  
date      | 1996-02-02  
host      | Golden State Warriors -- 114  
guest     | Toronto Raptors -- 111  
winner_drb | 14  
-[ RECORD 3 ]-----  
date      | 1998-03-31  
host      | Vancouver Grizzlies -- 101  
guest     | Dallas Mavericks -- 104  
winner_drb | 14  
-[ RECORD 4 ]-----  
date      | 2009-01-14  
host      | New York Knicks -- 128  
guest     | Washington Wizards -- 122  
winner_drb | 14
```

Pure SQL Histogram drawing



Pure SQL Histogram drawing

```
with drb_stats as (
    select min(drb) as min,
           max(drb) as max
      from team_stats
),
histogram as (
    select width_bucket(drb, min, max, 9) as bucket,
           int4range(min(drb), max(drb), '[]') as range,
           count(*) as freq
      from team_stats, drb_stats
     group by bucket
    order by bucket
)
select bucket, range, freq,
       repeat('*', (freq::float / max(freq) over() * 30)::int)
  from histogram;
```

Pure SQL Histogram drawing

bucket	range	freq	bar
1	[10,15)	52	
2	[15,20)	1363	**
3	[20,25)	8832	*****
4	[25,30)	20917	*****
5	[30,35)	20681	*****
6	[35,40)	9166	*****
7	[40,45)	2093	***
8	[45,50)	247	
9	[50,54)	20	
10	[54,55)	1	

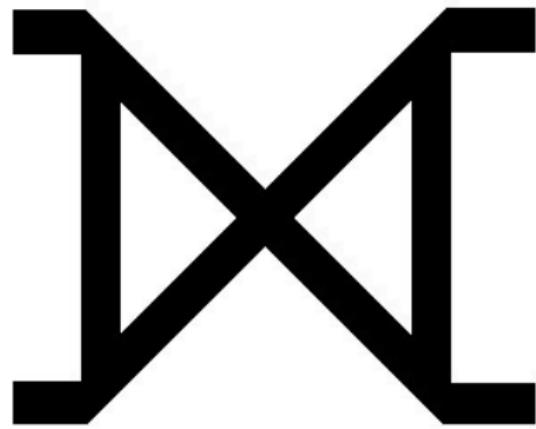
(10 rows)

Time: 53.570 ms

Writable CTE

```
with queue as (
    insert into queue (extension)
        select id
            from extension
            where shortname = $1
        returning id, extension
)
select q.id, e.id as ext_id,
    e.fullname, e.uri, e.description
from      queue q
    join extension e on q.extension = e.id;
```

PostgreSQL Joins



PostgreSQL Joins

```
WITH upd AS (
    UPDATE target t
        SET counter = t.counter + s.counter,
        FROM source s
        WHERE t.id = s.id
    RETURNING s.id
)
INSERT INTO target(id, counter)
    SELECT id, sum(counter)
        FROM source s LEFT JOIN upd t USING(id)
        WHERE t.id IS NULL
    GROUP BY s.id
    RETURNING t.id
```

PostgreSQL Lateral Joins

Fetch the 10 most recent news for each topic (**Top-N**)

```
SELECT n.*  
FROM subscriptions s  
JOIN LATERAL (SELECT *  
              FROM news n  
              WHERE n.topic = s.topic  
              ORDER BY n.created DESC  
              LIMIT 10  
            ) top_news ON (true)  
WHERE s.user_id = ?  
ORDER BY n.created DESC  
LIMIT 10
```



POSTGRESQL IS YESQL!



Questions?

Now is the time to ask!

