

MIGRATING TO POSTGRESQL, THE NEW STORY FOSDEM PGDAY, 2015

Dimitri Fontaine `dimitri@2ndQuadrant.fr`
`@tapoueh`

January 30, 2015




PRINCIPAL CONSULTANT AT 2NDQUADRANT



POSTGRESQL MAJOR CONTRIBUTOR


- pgloader
- prefix, skytools
- apt.postgresql.org
- CREATE EXTENSION
- CREATE EVENT TRIGGER
- *Bi-Directional Réplication*
- pginstall

PostgreSQL




POSTGRESQL IS YESQL!

PostgreSQL



Let's talk about MySQL for a minute

Just in the context of migrating from it, of course



Why Migrating from MySQL to PostgreSQL?

MySQL

- Storage Engine
- Single Application
- Data Loss with Replication
- Weak Data Types Validation
- Either transactions or
- Lack of

PostgreSQL

- Data Access Service
- Application Suite
- Durability and Availability
- Consistency
- Full Text Search, PostGIS
- Proper Documentation



The migration budget


What are the costs?

- Migrating the Data
- Migrating the Code
- Quality Assurance
- Opportunity Cost



The boring parts really are

MySQL used not to be so serious about data consistency... still is



JOHNSON AND ATKINSON'S TYPE CASTER.

Digitized by Google

Difficulties when migrating MySQL data

Dates and The *Gregorian* Calendar

```
MariaDB [talk]> create table dates(d datetime);
MariaDB [talk]> insert into dates
    values('0000-00-00'), ('0000-10-31'), ('2013-10-00');
```

```
MariaDB [talk]> select * from dates;
```

```
+-----+
| d                |
+-----+
| 0000-00-00 00:00:00 |
| 0000-10-31 00:00:00 |
| 2013-10-00 00:00:00 |
+-----+
3 rows in set (0.00 sec)
```

The GOD algorithm



The GOD algorithm

What you want to do instead

```
$ createdb pagila
$ pgloader mysql://user@localhost/sakila
    pgsql:///pagila
```



pgloader mysql://root@localhost/sakila pgsql:///pagila

./build/bin/pgloader mysql://root@localhost/sakila pgsql:///pagila

2015-01-30T12:49:09.030000+01:00 LOG Main logs in '/private/tmp/pgloader/pgloader.log'

2015-01-30T12:49:09.034000+01:00 LOG Data errors in '/private/tmp/pgloader/'

table name	read	imported	errors	time
fetch meta data	43	43	0	0.358s
create, drop	0	36	0	0.263s
actor	200	200	0	0.151s
address	603	603	0	0.067s
category	16	16	0	0.025s
city	600	600	0	0.036s
country	109	109	0	0.019s
customer	599	599	0	0.040s
film	1000	1000	0	0.106s
film_actor	5462	5462	0	0.171s
film_category	1000	1000	0	0.068s
film_text	1000	1000	0	0.070s
inventory	4581	4581	0	0.172s
language	6	6	0	0.078s
payment	16049	16049	0	0.470s
rental	16044	16044	0	0.458s
staff	2	2	0	0.103s
store	2	2	0	0.035s
Index Build Completion	0	0	0	0.001s
Create Indexes	41	41	0	0.291s
Reset Sequences	0	13	0	0.045s
Primary Keys	16	16	0	0.014s
Foreign Keys	22	22	0	0.079s
Comments	0	0	0	0.000s
Total import time	47273	47273	0	2.829s



pgloader: load data into PostgreSQL

`http://pgloader.io/`




pgloader: Open Source, github

`https://github.com/dimitri/pgloader`



pgloader: what about loading data?

`http://pgloader.io/`



pgloader main features

pgloader is built around copy

- Error handling and reject files
- On the fly data transformations
- Very simple command line for simple use cases
- Advanced command language for advanced use cases
- Parallelism to benefit from async IO
- Lots of input formats




`http://pgloader.io/howto/csv.html`



<http://pgloader.io/howto/quickstart.html>

PostgreSQL



dBase III

<http://pgloader.io/howto/dBase.html>



SQLite

<http://pgloader.io/howto/sqlite.html>



MySQL

`http://pgloader.io/howto/mysql.html`



<http://pgloader.io/>



Microsoft®
SQL Server®

And more to come

File formats with on-the-fly normalisation

<?xml?>

{JSON}

Other database systems

ORACLE®

 INFORMIX

 Microsoft®
SQL Server®

 SYBASE®

pgloader database migration process

The data migration process, step by step

- 1 Fetch metadata from source database catalogs
- 2 Prepare PostgreSQL database
- 3 COPY data
- 4 Complete PostgreSQL database
- 5 Display summary (human readable, json, csv, copy)



Fetching Metadata

Currently supported metadata

- Schemas
- Tables
- *Columns*
- Indexes
- Constraints
- Comments (mysql only)
- *Materializing Views* (mysql only)



Prepare PostgreSQL database

Prepare PostgreSQL for receiving the data

- Schemas
- Tables
- Columns
- Rename indexes with table oids in memory



Copy Data

Copy the data from the source to the target

- For each table, COPY data in
- Two threads work in parallel (reader/writer)
- Then for each table, create all indexes in parallel
- Currently no cap on the number of threads



Complete PostgreSQL Database


Install constraints

- Reset **Sequences**
- Upgrade unique indexes into **Primary Keys** where required
- **Foreign Keys**
- **Comments**



Migrating columns

What do you mean **columns**?



JOHNSON AND ATKINSON'S TYPE CASTER.

Digitized by Google

A simple use case

Remember that example? we'll see a more detailed one...

```
$ createdb pagila
$ pgloader mysql://user@localhost/sakila
    pgsql:///pagila

$ pgloader migration.load
```



An advanced use case 1/3

```
LOAD DATABASE
```

```
FROM      mysql://root@unix:/tmp/mysql.sock:/goeuro  
INTO postgresql://dim@unix:/tmp:/godollar
```

```
-- INCLUDING ONLY TABLE NAMES MATCHING ~/sizes/  
EXCLUDING TABLE NAMES MATCHING ~/encoding/
```

```
DECODING TABLE NAMES
```

```
    MATCHING ~/messed/,  
            ~/encoding/  
    AS utf-8
```

```
SET maintenance_work_mem to '1 GB'
```



An advanced use case 2/3

```
CAST datetime to timestampz
      drop default drop not null
      using zero-dates-to-null,

column bools.a to boolean drop typemod
      using tinyint-to-boolean,

type char when (= precision 1)
      to char keep typemod,

column ascii.s using byte-vector-to-bytea,

column enumerate.foo
      using empty-string-to-null
```



An advanced use case 3/3

```
WITH lowercase identifiers,  
    truncate,  
    -- data only,  
    -- schema only,  
    create tables,  
    include drop,  
    create indexes,  
    reset sequences,  
    foreign keys;  
  
-- quote identifiers  
  
-- create no tables  
-- include no drop  
-- create no indexes  
-- reset no sequences  
-- no foreign keys
```



pgloader: load data into PostgreSQL

`http://pgloader.io/`



Questions?

Now is the time to ask!

