


PostgreSQL Extensions

Advanced Use Cases

Dimitri Fontaine `dimitri@2ndQuadrant.fr`

Mercredi 20 Mai 2015



PRINCIPAL CONSULTANT AT
2NDQUADRANT



PostgreSQL Extensions

```
CREATE EXTENSION ...;
```

Advanced Extension Use Cases

Agenda

- How PostgreSQL extensibility works
- Things you can do with a PostgreSQL Extension
- The PostgreSQL indexing Framework
- **How to solve some practical use cases with existing extensions**



PostgreSQL Extensibility

PostgreSQL is *highly* **extensible**



PostgreSQL Extensibility

```
select col1, col2 from table where col1 = 'something';
```



PostgreSQL Extensibility

```
SELECT col  
  FROM table  
 WHERE stamped > date 'today' - interval '1 day';
```



PostgreSQL Extensibility

```
select iprange, locid
  from geolite.blocks
 where iprange >= '91.121.37.122';
```

iprange	locid
91.121.0.0-91.121.159.255	75

(1 row)

Time: 1.220 ms



PostgreSQL Extensibility: Operator Classes

SQL Operators are all dynamic and found in the catalogs

```
select amopopr::regoperator
  from pg_opclass c
 join pg_am am
    on am.oid = c.opcmethod
 join pg_amop amop
    on amop.amopfamily = c.opcfamily
 where opcintype = 'ip4r'::regtype
    and am.amname = 'gist';
```

```
amopopr
-----
>>=(ip4r,ip4r)
<<=(ip4r,ip4r)
>>(ip4r,ip4r)
<<(ip4r,ip4r)
&&(ip4r,ip4r)
=(ip4r,ip4r)
(6 rows)
```



PostgreSQL Extensibility

```
select id, name, pos,  
       pos <@> point(-6.25,53.34) as miles  
from pubnames  
order by pos <-> point(-6.25,53.34)  
limit 10;
```



PostgreSQL is Extensible

PostgreSQL plugins are data types and index support


- Data Type
- Input/Output functions
- Casts
- Operator Classes



PostgreSQL is Extensible

PostgreSQL support several kind of indexes


- BTree, binary tree
- GiST, Generalized Search Tree
- SP-GiST, Space Partitioned GiST
- GIN, Generalized Inverted Index



Binary Tree

Btree, the default index type

- Built for speed
- *unique* concurrency tricks
- Balanced
- support function: `cmp`
- operators: `<=` `<` `=` `>` `>=`




Generalized Index Search Tree

GiST or the Indexing API

- Built for comfort
- Balanced
- API: consistent, same, union
- API: penalty, picksplit
- API: compress, decompress
- operators: $@>$ $<@$ $\&\&$ $@@$ $=$ $\&<$ $\&>$
 $<<|$...


R-tree Hierarchy




Generalized Inverted iNdex

Indexing several pointers per value, inversed cardinality

- Built for Text Search and Arrays
- Balanced
- API: compare, consistent
- API: extractValue, extractQuery
- operators: @> <@ && =



Extensions and data types




Some extensions example

46 Contribs, Community extensions, Private ones...

- **hll**
- **cube**
- **ltree**
- **citext**
- **hstore**
- **earthdistance**
- **pgq**
- **pg_trgm**
- **wildspeed**
- **plproxy**
- **PostGIS**
- **ip4r**
- **intarray**
- **prefix**
- **pgfincore**
- **pgcrypto**
- **pgstattuple**
- **pg_buffercache**
- **pg_stat_statements**
- **pgfincore**



IP Ranges, ip4r



IP Ranges, ip4r

```
table geolite.blocks limit 10;
```

iprange	locid
-----+-----	
1.0.0.0/24	17
1.0.1.0-1.0.3.255	49
1.0.4.0/23	14409
1.0.6.0/23	17
1.0.8.0/21	49
1.0.16.0/20	14614
1.0.32.0/19	47667
1.0.64.0/18	111
1.0.128.0-1.0.147.255	209
1.0.148.0/24	22537
(10 rows)	



IP Ranges, ip4r, Geolocation

PostgreSQL allows using SQL and JOINS to match IP4R with geolocation.

```
select *  
  from geolite.blocks  
 join geolite.location  
   using(locid)  
 where iprange  
        >>=  
        '74.125.195.147';
```



IP Ranges, ip4r, Geolocation

PostgreSQL allows using SQL and JOINS to match IP4R with geolocation.


```
select *
  from geolite.blocks
 join geolite.location
    using(loid)
 where iprange
        >=
        '74.125.195.147';
```

-[RECORD 1]-----	
loid	2703
iprange	74.125.189.24-74.125.
country	US
region	CA
city	Mountain View
postalcode	94043
location	(-122.0574,37.4192)
metrocode	807
areacode	650

Time: 1.335 ms




Earth Distance



How Far is The Nearest Pub

The point datatype is in-core

```
# CREATE TABLE pubnames
(  
  id      bigint,  
  pos     POINT,  
  name    text  
);
```



How Far is The Nearest Pub

```
select name, pos
  from pubnames
order by pos <-> point(-6.25,53.346)
 limit 3;
```

Pub Name	pos
Ned's	(-6.2519967,53.3458267)
Sub Lounge	(-6.2542332,53.3469085)
O'Neill's of Pearse Street	(-6.2524389,53.3448589)

(3 rows)

Time: 18.679 ms



How Far is The Nearest Pub

```
CREATE INDEX on pubnames USING GIST(pos);
```

```
select name,  
       pos  
  from pubnames  
 order by pos <-> point(-0.12,51)  
  limit 3;
```

name	pos
Ned's	(-6.25,53.34
Sub Lo	(-6.25,53.34
O'Neil	(-6.25,53.34

(3 rows)

Time: 0.849 ms



How Far is The Nearest Pub, in Miles please.

```
create extension cube;  
create extension earthdistance;
```

```
select name,  
       pos <@> point(-6.25,53.34) miles  
       from pubnames  
order by pos <-> point(-6.25,53.34)  
limit 3;
```

name		miles
Ned's		0.06
Sub Lo		0.07
O'Neil		0.12

(3 rows)

Time: 1.335 ms



Some pubs far away from here...

```
select c.name as city,  
pos <@> point(-6.25,53.34) as miles  
from pubnames p,  
    lateral (select name  
              from cities c  
              order by c.pos <-> p.pos  
              limit 1) c  
order by pos <-> point(-6.25,53.34)  
        desc  
limit 5;
```

city		miles
-----	+	-----
Canterbury		399.44
Canterbury		378.91
Canterbury		392.08
Canterbury		397.30
Canterbury		379.68

(5 rows)

Time: 636.445 ms



Geolocation: ip4r meets earthdistance



Some pubs nearby... some place...

```
with geoloc as
(
  select location as l
    from location
   join blocks using(loid)
  where iprange
        >>=
        '212.58.251.195'
)
select name,
       pos <@> 1 miles
  from pubnames, geoloc
order by pos <-> 1
limit 10;
```

name	miles
Blue Anchor	0.299
Dukes Head	0.360
Blue Ball	0.337
Bell (aka The Rat)	0.481
on the Green	0.602
Fox & Hounds	0.549
Chequers	0.712
Sportsman	1.377
Kingswood Arms	1.205
Tattenham Corner	2.007

(10 rows)

Time: 3.275 ms



Trigrams



Trigrams and similarity

similar but not quite *like* the same

```
create extension pg_trgm;
```

```
select show_trgm('tomy') as tomy,  
       show_trgm('Tomy') as "Tomy",  
       show_trgm('tom torn') as "tom torn",  
       similarity('tomy', 'tom'),  
       similarity('dim', 'tom');
```

```
-[ RECORD 1 ]-----  
tomy          | {" t", " to", "my ",omy,tom}  
Tomy          | {" t", " to", "my ",omy,tom}  
tom torn      | {" t", " to", "om ",orn,"rn ",tom,tor}  
similarity    | 0.5  
similarity    | 0
```



Trigrams and typos

Use your data to help your users out

```
select actor
  from products
 where actor ~* 'tomy';
actor
```

(0 rows)

Time: <unregistered>

```
select actor
  from products
 where actor % 'tomy';
actor
```

TOM TORN

TOM DAY


(2 rows)

Time: 26.972 ms



Trigrams search indexing

```
create index on products using gist(actor gist_trgm_ops);
```



```
select actor  
  from products  
 where actor % 'tomy';  
      actor
```

```
-----  
TOM TORN  
TOM DAY  
(2 rows)
```

```
Time: 2.695 ms
```



Trigrams and autocompletion

Use your data to help your users out

```
explain (costs off)
  select * from products where actor ~* 'tomy';
          QUERY PLAN
```

```
Index Scan using products_actor_idx on products
  Index Cond: ((actor)::text ~* 'tomy'::text)
(2 rows)
```



Trigrams and autocompletion

Use your data to help your users out

```
select actor
  from products
 where actor % 'fran'
order by actor <-> 'fran'
 limit 10;
```

actor


```
FRANK HAWKE
FRANK BERRY
FRANK POSEY
FRANK HAWKE
FRANCES DEE
FRANK LEIGH
FRANCES DAY
FRANK FOSTER
FRANK HORNE
FRANK TOMEI
(10 rows)
```

Time: 2.960 ms





PL/Proxy is all about *Sharding*



We're going to use it for *Remote Procedure Call*



Classic Auditing

```
create table example
(
    id serial,
    f1 text,
    f2 text
);
```

```
create table audit
(
    change_date timestampz
                default now(),
    before hstore,
    after  hstore
);
```



Classic trigger based Auditing

```
~# begin;
~*# update example set f1 = 'b' where id = 1;
~*# rollback;

~# select * from audit;
  change_date | before | after
-----+-----+-----
(0 rows)
```



Setting up PL/Proxy

```
create extension plproxy;

create server local
foreign data wrapper plproxy
options (p0 'dbname=dim');

create user mapping
for public
server local
options (user 'dim');
```



PL/Proxy: Basic Testing

```
create function test_proxy
    (i int)
    returns int
    language plproxy
as '
    cluster ''local'';
    select i;
';
```

```
select test_proxy(1);
test_proxy
-----
          1
(1 row)

Time: 0.866 ms
```



Implementing Autonomous Transactions for Auditing

REMOTE TRIGGER



Trigger Functions 1/3: the trigger

```
create function audit_trigger()  
  returns trigger  
  language plpgsql  
as '  
begin  
  perform audit_proxy(old, new);  
  return new;  
end;  
';
```



Trigger Functions 2/3: the proxy

```
create function audit_proxy
(
    old example,
    new example
)
returns void
language plproxy
as '
    cluster ''local'';
    target audit;
';
```



Trigger Functions 3/3: the implementation

```
create function audit
(
    old example,
    new example
)
returns void
language SQL
as '
    INSERT INTO audit(before, after)
        SELECT hstore(old), hstore(new);
';
```



Trigger Definition

```
drop trigger if exists audit on example;

create trigger audit
  after update on example
    for each row
      -- careful, defaults to FOR EACH STATEMENT!
execute procedure audit_trigger();
```



Autonomous Auditing Transaction

```
~# begin;
```

```
BEGIN
```

```
~*# update example set f1 = 'b' where id = 1;
```

```
UPDATE 1
```

```
~*# rollback;
```

```
ROLLBACK
```


Autonomous Auditing Transaction

We did ROLLBACK; the transaction

```
~# select change_date,  
         before,  
         after,  
         after - before as diff  
       from audit;  
-[ RECORD 1 ]-----  
change_date | 2013-10-14 14:29:09.685105+02  
before      | "f1"=>"a", "f2"=>"a", "id"=>"1"  
after       | "f1"=>"b", "f2"=>"a", "id"=>"1"  
diff        | "f1"=>"b"
```



State of The Art Cardinality Estimation Algorithm



Creating the unique visitors tracking table

```
CREATE EXTENSION hll;

-- Create the destination table
CREATE TABLE daily_uniques (
    DATE            DATE UNIQUE,
    users           hll
);

-- Our first aggregate update
UPDATE daily_uniques
    SET users = hll_add(users,
                        hll_hash_text('123.123.123.123'))
WHERE date = current_date;
```



Production ready updates

```
--  
-- First upload a new batch, e.g. using  
-- CREATE TEMP TABLE new_batch as VALUES(), (), ...;  
--  
WITH hll(agg) AS (  
    SELECT hll_add_agg(hll_hash_text(value))  
        FROM new_batch  
)  
UPDATE daily_uniques  
    SET users = CASE WHEN hll.agg IS NULL THEN users  
                    ELSE hll_union(users, hll.agg)  
                END  
FROM hll  
WHERE date = current_date;
```



Daily Reporting

```
with stats as (  
    select date,  
           #users as daily,  
  
           #hll_union_agg(users)  
           over() as total  
    from daily_uniques  
)  
select date,  
       daily,  
       daily/total*100  
    from stats  
order by date;
```

date	daily	percent
2013-02-22	401677	25.19
2013-02-23	660187	41.41
2013-02-24	869980	54.56
2013-02-25	154996	9.72

(4 rows)



Monthly Reporting

```
select to_char(date, 'YYYY/MM'),  
       #hll_union_agg(users)  
from daily_uniques  
group by 1;
```

month	monthly
2013/02	1960380

(1 row)



PostgreSQL is YeSQL!



Recap

We saw a number of extensions, each with a practical use case

ip4r IP Ranges and Geolocation

Earth Longitude, Latitude, Computing distances on a map

Trigrams Fixing typos, autocompletion

PL/Proxy Sharding, RPC, Autonomous Transactions

HLL Cardinalities, Unique Visitors



Questions?

Now is the time to ask!

