

UTILISER POSTGRESQL EN 2015

Java User Group, Montpellier

Dimitri Fontaine dimitri@2ndQuadrant.fr
@tapoueh

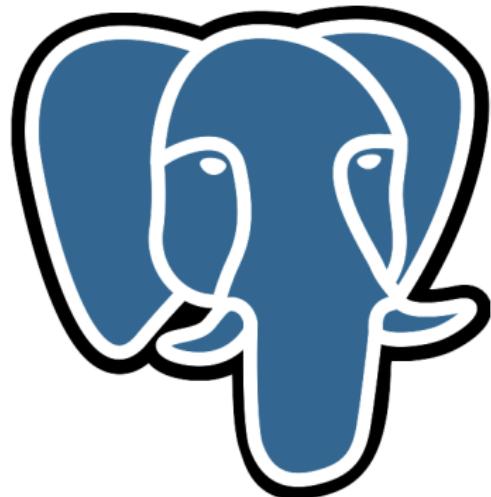
20 Mai 2015

PRINCIPAL CONSULTANT AT 2NDQUADRANT



POSTGRESQL MAJOR CONTRIBUTOR

- CREATE EXTENSION
- CREATE EVENT TRIGGER
- *Bi-Directional Réplication*
- pgloader
- prefix, skytools
- apt.postgresql.org
- pginstall
- pgcharts



<http://pgloader.io/>



Migrer de MySQL à PostgreSQL

Example de migration des données, *complet*

```
$ createdb pagila
$ pgloader mysql://user@localhost/sakila
                  pgsql:////pagila
```

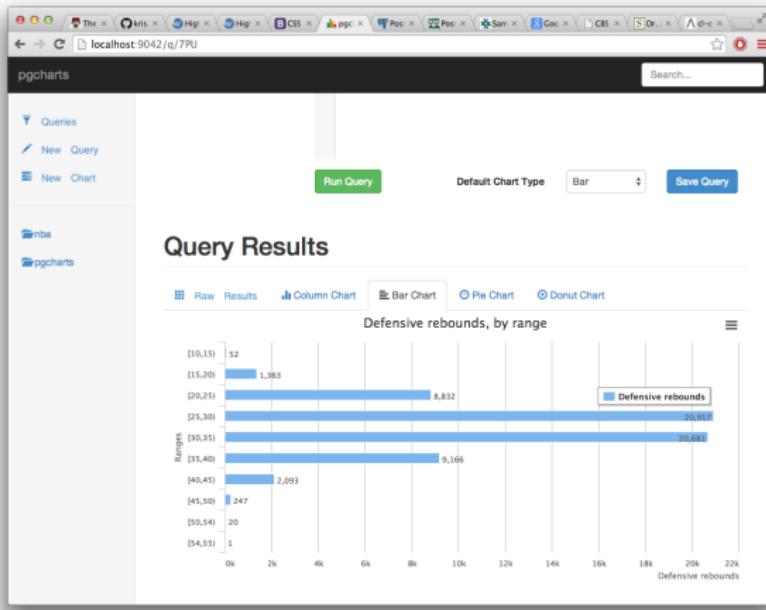
```
pgloader mysql://root@localhost/sakila pgsql:///pagila
```

```
./build/bin/pgloader mysql://root@localhost/sakila pgsql:///pagila
2015-01-30T12:49:09.030000+01:00 LOG Main logs in '/private/tmp/pgloader/pgloader.log'
2015-01-30T12:49:09.034000+01:00 LOG Data errors in '/private/tmp/pgloader/'
```

table name	read	imported	errors	time
fetch meta data	43	43	0	0.358s
create, drop	0	36	0	0.263s
-----	-----	-----	-----	-----
actor	200	200	0	0.151s
address	603	603	0	0.067s
category	16	16	0	0.025s
city	600	600	0	0.036s
country	109	109	0	0.019s
customer	599	599	0	0.040s
film	1000	1000	0	0.106s
film_actor	5462	5462	0	0.171s
film_category	1000	1000	0	0.068s
film_text	1000	1000	0	0.070s
inventory	4581	4581	0	0.172s
language	6	6	0	0.078s
payment	16049	16049	0	0.470s
rental	16044	16044	0	0.458s
staff	2	2	0	0.103s
store	2	2	0	0.035s
Index Build Completion	0	0	0	0.001s
-----	-----	-----	-----	-----
Create Indexes	41	41	0	0.291s
Reset Sequences	0	13	0	0.045s
Primary Keys	16	16	0	0.014s
Foreign Keys	22	22	0	0.079s
Comments	0	0	0	0.000s
-----	-----	-----	-----	-----
Total import time	47273	47273	0	2.829s

Turn your SQL queries into nice looking charts.

<https://github.com/dimitri/pgcharts/>

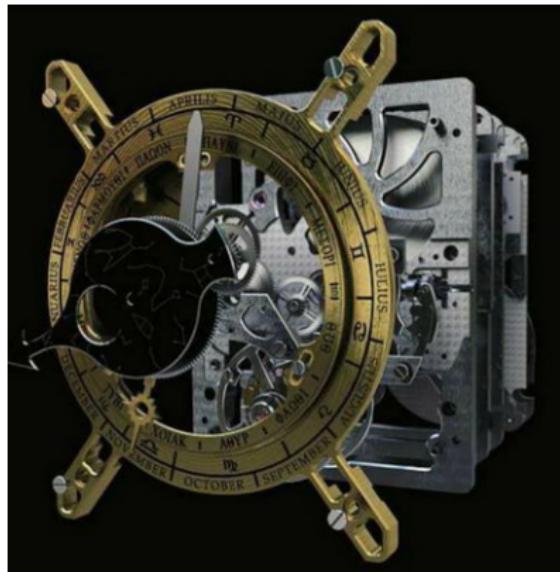


Système de Gestion de Bases de données Relationnelles



La Fabuleuse machine d'Anticythère

Système de Gestion de Bases de données Relationnelles



La Fabuleuse machine d'Anticythère

Système de Gestion de Bases de données Relationnelles





Rollback

Transactions

```
BEGIN;  
  DROP TABLE donnees_critiques;  
ROLLBACK;
```

Transactions

```
START TRANSACTION ISOLATION LEVEL SERIALIZABLE;  
DROP TABLE donnees_critiques;  
ROLLBACK;
```

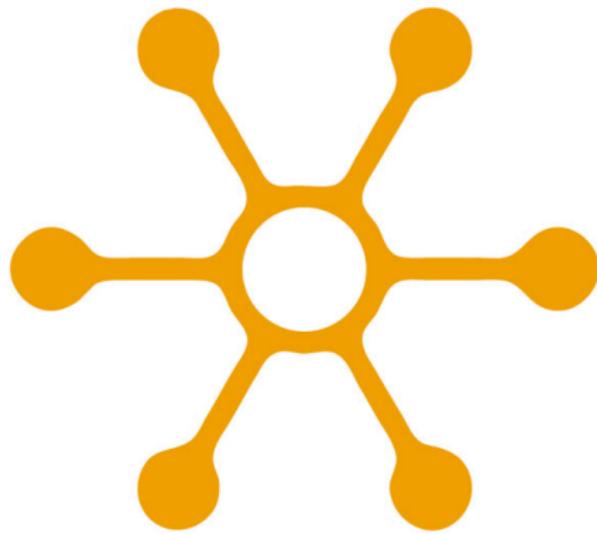
Transactions

```
BEGIN;  
    CREATE INDEX ... ON ...;  
    EXPLAIN ANALYZE SELECT ...;  
ROLLBACK;
```

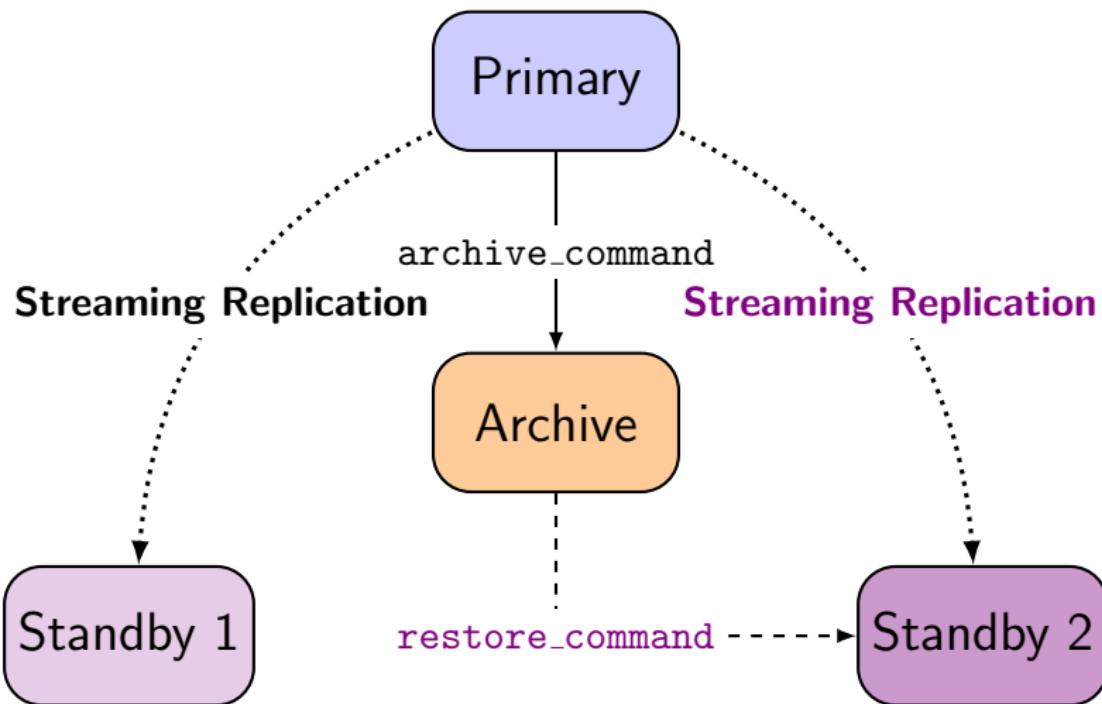
Backups cohérents à chaud



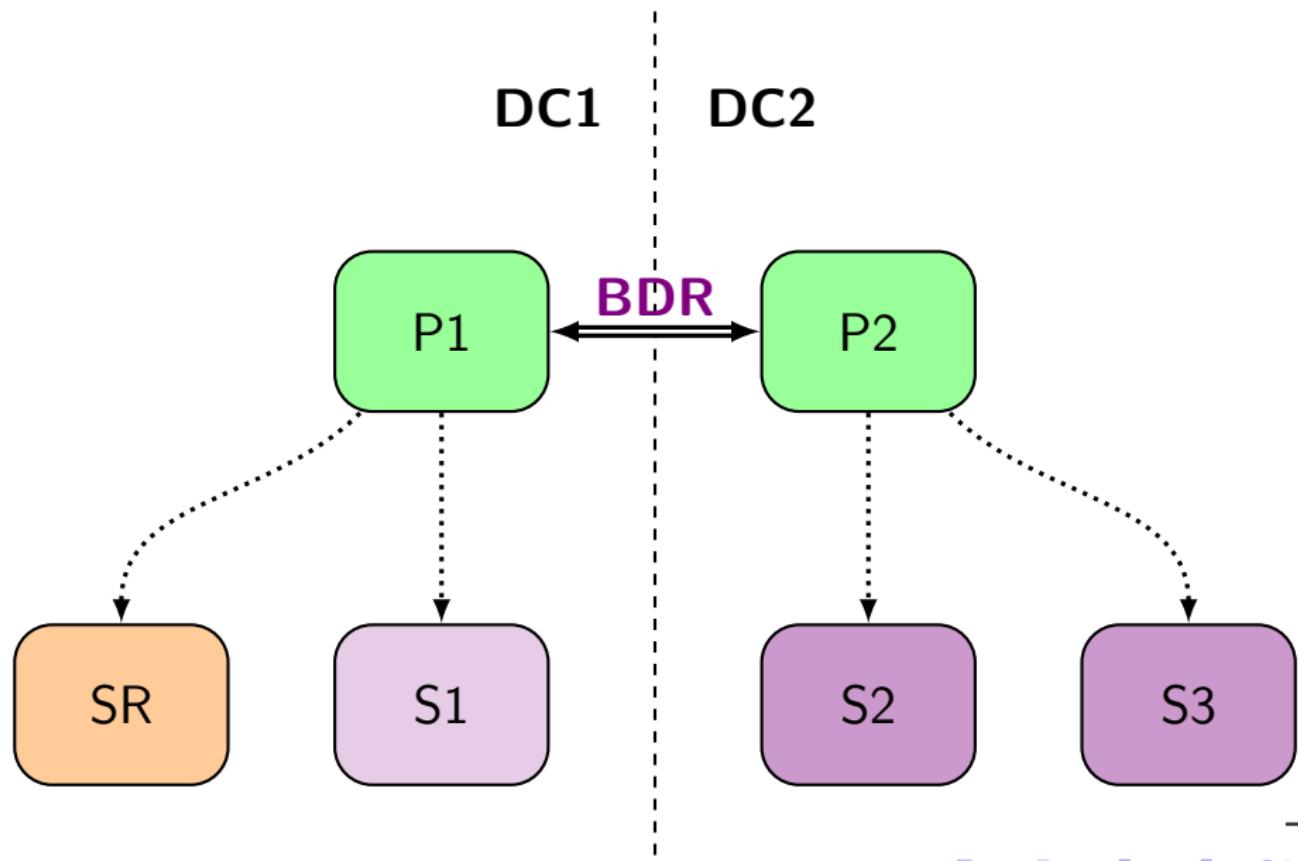
Disponibilité des services et des données



Disponibilité des services et des données



Disponibilité des services et des données



Cohérence des données



Type de données et contraintes

Types et Intégrité des données

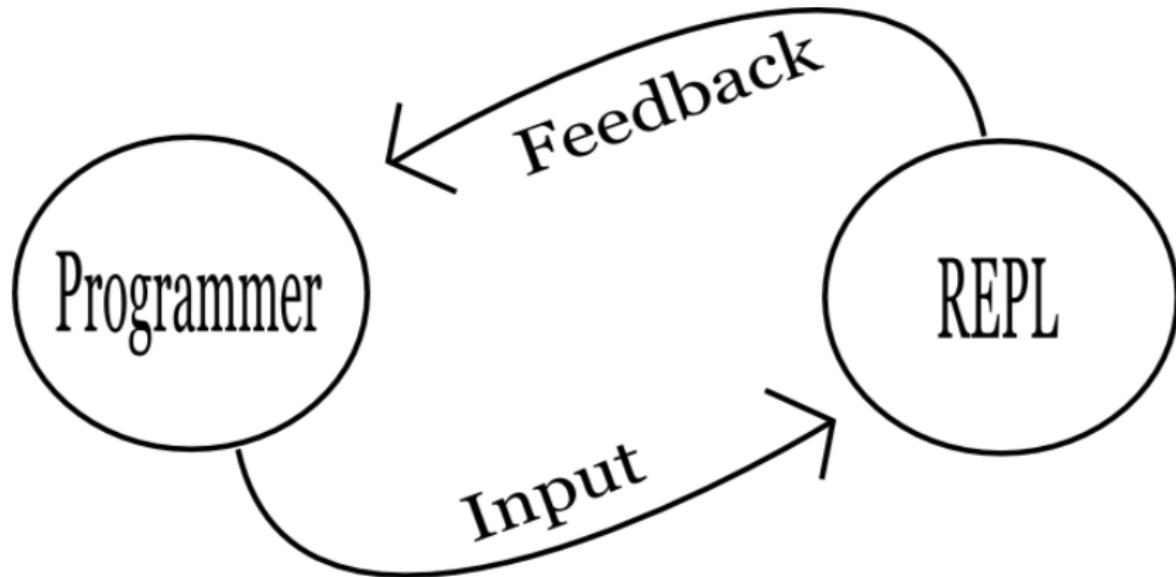
- Integer
- Arbitrary precision numbers, UUID
- Floating point
- Character, Text
- Bytea, bitstring
- Date/Time, Time Zones
- Boolean
- Enum, Arrays, Composite Types, Range Types
- Point, Line Segments, Boxes
- Paths, Polygons, Circles
- Inet, CIDR, Macaddr
- JSON, XML

Type de données et contraintes

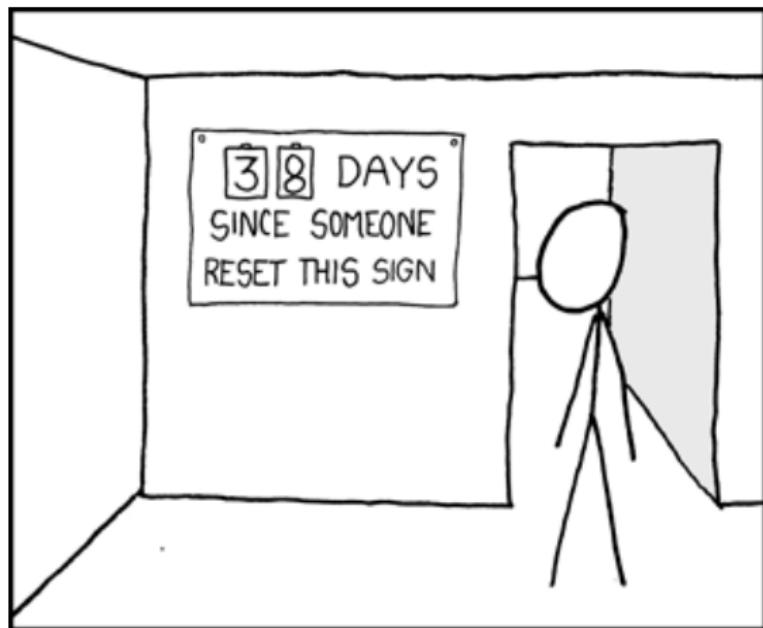
```
CREATE TABLE reservation
(
    room      text,
    professor text,
    during    period,
    EXCLUDE USING gist
    (
        room with =,
        during with &&
    )
);
```

```
CREATE TABLE circles (
    c circle,
    EXCLUDE USING gist (c WITH &&)
);
```

La console interactive psql



A simple project



SQL: we start with DDLs

Joe Celko: 80% of the job is to define the schema

Example (DDL)

```
create table measures(date timestamptz primary key,  
                      mesure integer);
```

```
dim=# \d measures
```

```
\d measures
```

Table "public.measures"

Column	Type	Modifiers
date	timestamp with time zone	not null
mesure	integer	

Indexes:

"measures_pkey" PRIMARY KEY, btree (date)

We take a very simple model for the presentation

```
create table measures(tick int, nb int);  
  
insert into measures  
values (1, 0), (2, 10), (3, 20), (4, 30), (5, 40),  
(6, 0), (7, 20), (8, 30), (9, 60);
```

Testing data

Let's take some measures as if they came out of our counter, starting at 0, and with a *reset* in there. In that example, the global usage measured is $40 + 60 = 100$.

```
select * from measures;
   tick | nb
-----+-----
      1 |  0
      2 | 10
      3 | 20
      4 | 30
      5 | 40
      6 |  0
      7 | 20
      8 | 30
      9 | 60
(9 rows)
```

Aside: PostgreSQL knows about arrays

```
select array_agg(nb) from measures;  
array_agg
```

```
-----  
{0,10,20,30,40,0,20,30,60}  
(1 row)
```

Finding the last counter value before reset

Write some SQL here

tick	nb	max
1	0	
2	10	
3	20	
4	30	
5	40	40
6	0	
7	20	
8	30	
9	60	60

(9 rows)

Window Functions: lead() over()

```
select tick,
       nb,
       lead(nb) over (order by tick)
  from measures;
```

tick	nb	lead
1	0	10
2	10	20
3	20	30
4	30	40
5	40	0
6	0	20
7	20	30
8	30	60
9	60	

(9 rows)

Window Functions and CASE

```
select tick, nb,
       case when lead(nb) over w < nb
             then nb
             when lead(nb) over w is null
             then nb
             else null
         end as max
  from measures
 window w as (order by tick);
```

tick	nb	max
1	0	
2	10	
3	20	
4	30	
5	40	40
6	0	
7	20	
8	30	
9	60	60

(9 rows)

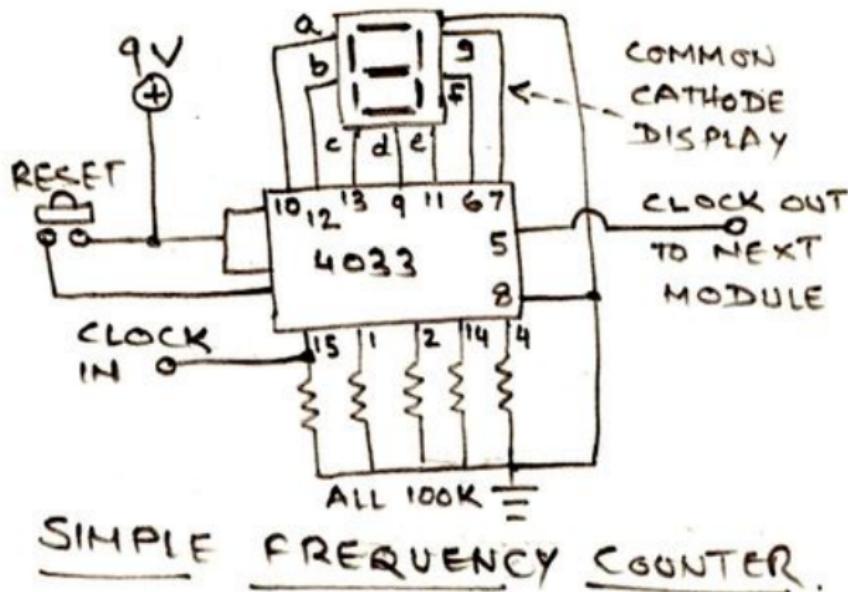
Window Functions and WHERE clause

```
with t(tick, nb, max) as (
    select tick, nb,
        case when lead(nb) over w < nb then nb
              when lead(nb) over w is null then nb
              else null
        end as max
    from measures
    window w as (order by tick)
)
select tick, nb, max from t where max is not null;
tick | nb | max
-----+----+-----
 5 | 40 | 40
 9 | 60 | 60
(2 rows)
```

Common Table Expressions to complement WITH

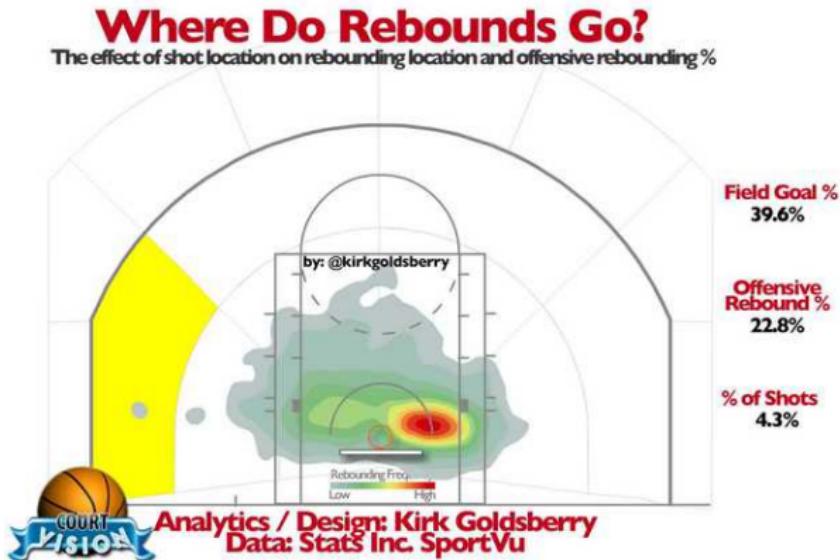
```
with t(tops) as (
    select case when lead(nb) over w < nb then nb
                when lead(nb) over w is null then nb
                else null
            end as max
      from measures
     window w as (order by tick)
)
select sum(tops) from t;
sum
-----
100
(1 row)
```

Getting usage from the counter: done. SQL. 9 lines.



STRUCTURED QUERY LANGUAGE





NBA Games Statistics

An interesting factoid: the team that recorded the fewest defensive rebounds in a win was the 1995-96 Toronto Raptors, who beat the Milwaukee Bucks 93-87 on 12/26/1995 despite recording only 14 defensive rebounds.

PostgreSQL Data Mining

```
with stats(game, team, drb, min) as (
    select ts.game, ts.team, drb, min(drb) over ()
        from team_stats ts
            join winners w on w.id = ts.game
                and w.winner = ts.team
)
select game.date::date,
    host.name || ' -- ' || host_score as host,
    guest.name || ' -- ' || guest_score as guest,
    stats.drb as winner_drb
from stats
    join game on game.id = stats.game
    join team host on host.id = game.host
    join team guest on guest.id = game.guest
where drb = min;
```

PostgreSQL Data Mining

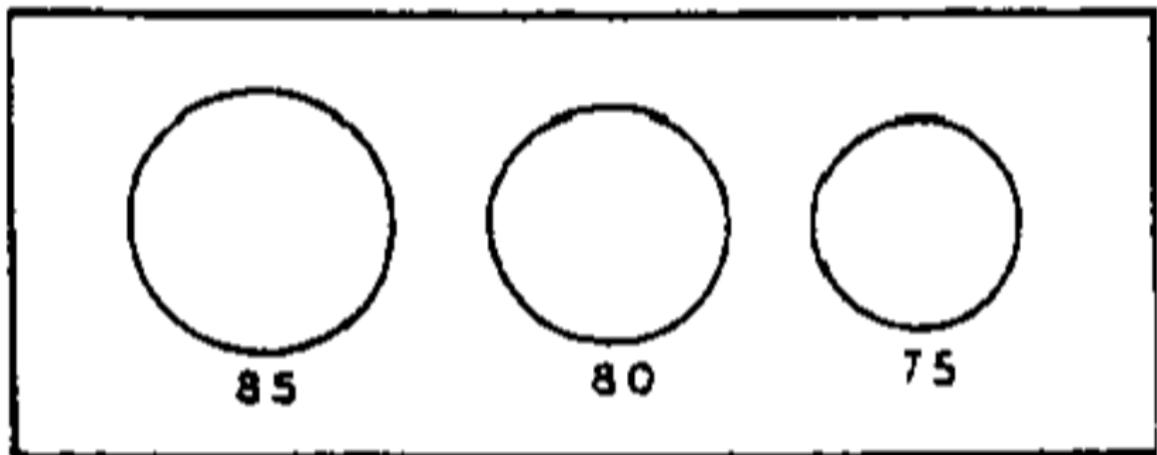
```
- [ RECORD 1 ] -----
date          | 1995-12-26
host          | Toronto Raptors -- 93
guest         | Milwaukee Bucks -- 87
winner_drb    | 14

-[ RECORD 2 ] -----
date          | 1996-02-02
host          | Golden State Warriors -- 114
guest         | Toronto Raptors -- 111
winner_drb    | 14

-[ RECORD 3 ] -----
date          | 1998-03-31
host          | Vancouver Grizzlies -- 101
guest         | Dallas Mavericks -- 104
winner_drb    | 14

-[ RECORD 4 ] -----
date          | 2009-01-14
```

Un histogramme tout en SQL



Histogramme tout en SQL

```
with drb_stats as (
    select min(drb) as min,
           max(drb) as max
      from team_stats
),
      histogram as (
        select width_bucket(drb, min, max, 9) as bucket,
               int4range(min(drb), max(drb), '[]') as range,
               count(*) as freq
      from team_stats, drb_stats
     group by bucket
    order by bucket
)
select bucket, range, freq,
       repeat('*', (freq::float / max(freq) over() * 30)::int)
  from histogram;
```

Histogramme tout en SQL

bucket	range	freq	bar
1	[10,15)	52	
2	[15,20)	1363	**
3	[20,25)	8832	*****
4	[25,30)	20917	*****
5	[30,35)	20681	*****
6	[35,40)	9166	*****
7	[40,45)	2093	***
8	[45,50)	247	
9	[50,54)	20	
10	[54,55)	1	

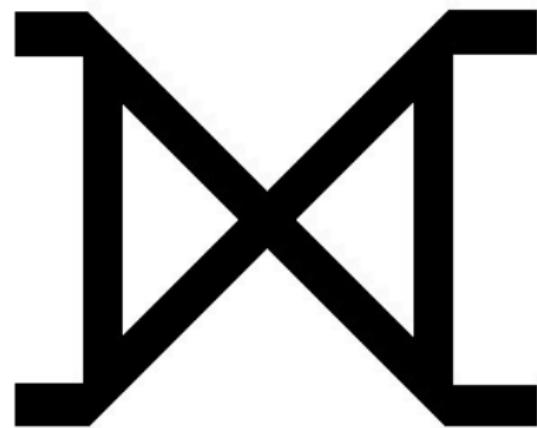
(10 rows)

Time: 53.570 ms

Savoir contourner les ORMs avec les wCTE

```
with queue as (
    insert into queue (extension)
        select id
            from extension
            where shortname = $1
        returning id, extension
)
select q.id, e.id as ext_id,
    e.fullname, e.uri, e.description
from      queue q
    join extension e on q.extension = e.id;
```

Les jointures avec PostgreSQL

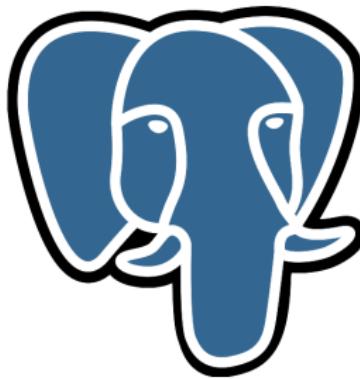


Les jointures avec PostgreSQL

```
WITH upd AS (
    UPDATE target t
        SET counter = t.counter + s.counter,
        FROM source s
        WHERE t.id = s.id
    RETURNING s.id
)
INSERT INTO target(id, counter)
    SELECT id, sum(counter)
        FROM source s LEFT JOIN upd t USING(id)
        WHERE t.id IS NULL
    GROUP BY s.id
    RETURNING t.id
```



POSTGRESQL IS YESQL!



Questions?

Now is the time to ask!

