

# UTILISER POSTGRESQL EN 2014

## PHP Tour, 2014

*Dimitri Fontaine* dimitri@2ndQuadrant.fr  
@tapoueh

23 juin 2014


PRINCIPAL CONSULTANT AT 2NDQUADRANT



## POSTGRESQL MAJOR CONTRIBUTOR

- pgloader
- prefix, skytools
- apt.postgresql.org
- CREATE EXTENSION
- CREATE EVENT TRIGGER
- *Bi-Directional Réplication*
- pginstall

PostgreSQL




# Système de Gestion de Bases de données Relationnelles



*La Fabuleuse machine d'Anticythère*

# Système de Gestion de Bases de données Relationnelles



*La Fabuleuse machine d'Anticythère*

# Système de Gestion de Bases de données Relationnelles





# Rollback

# Transactions

```
BEGIN;  
  DROP TABLE donnees_critiques;  
ROLLBACK;
```

# Transactions

```
START TRANSACTION ISOLATION LEVEL SERIALIZABLE;  
DROP TABLE donnees_critiques;  
ROLLBACK;
```


# Transactions

```
BEGIN;  
    CREATE INDEX ... ON ...;  
    EXPLAIN ANALYZE SELECT ...;  
ROLLBACK;
```


# Backups cohérents à chaud



# Disponibilité des services et des données



# Disponibilité des services et des données



# Cohérence des données




# Type de données et contraintes

```
CREATE TABLE reservation (
    during tsrange NOT NULL,
    EXCLUDE USING gist (during WITH &&)
);
```

```
CREATE TABLE circles (
    c circle,
    EXCLUDE USING gist (c WITH &&)
);
```

# La console interactive psql



# Types de données avancés

```
select iprange, locid
  from geolite.blocks
 where iprange >= '91.121.37.122';
```

iprange		locid
91.121.0.0-91.121.159.255		75

(1 row)

Time: 1.220 ms

# IP Ranges, ip4r, Geolocation

Géolocalisation et métadonnées obtenues sur une jointure

```
select *  
from geolite.blocks  
join geolite.location  
using(locid)  
where iprange  
      >>=  
      '74.125.195.147';
```

```
- [ RECORD 1 ] -----  
locid          | 2703  
iprange        | 74.125.189.24-74.125.  
country        | US  
region         | CA  
city           | Mountain View  
postalcode     | 94043  
location        | (-122.0574,37.4192)  
metrocode      | 807  
areacode       | 650
```

Time: 1.335 ms


# STRUCTURED QUERY LANGUAGE




# How Far is The Nearest Pub

The point datatype is in-core

```
CREATE TABLE pubnames
(
    id      bigint,
    pos     POINT,
    name    text
);
```



# Earth Distance



# How Far is The Nearest Pub, in Miles please.

```
create extension cube;  
create extension earthdistance;
```

```
select name,  
       pos <@> point(-6.25,53.34) miles  
  from pubnames  
order by pos <-> point(-6.25,53.34)  
   limit 3;
```

name	miles
Ned's	0.06
Sub Lo	0.07
O'Neil	0.12

(3 rows)

Time: 1.335 ms

# Geolocation



# Geolocalisation et meta-données

```
with geoloc as
(
  select location as l
  from location
  join blocks using(locid)
  where iprange
    >>=
    '212.58.251.195'
)
select name,
       pos <@> 1 miles
  from pubnames, geoloc
 order by pos <-> 1
 limit 10;
```

name	miles
Blue Anchor	0.299
Dukes Head	0.360
Blue Ball	0.337
Bell (aka The Rat)	0.481
on the Green	0.602
Fox & Hounds	0.549
Chequers	0.712
Sportsman	1.377
Kingswood Arms	1.205
Tottenham Corner	2.007
(10 rows)	

Time: 3.275 ms

# Trigrams



# Trigrammes et autocompletion

```
select actor  
      from products  
     where actor % 'fran'  
order by actor <-> 'fran'  
     limit 10;
```

actor
FRANK HAWKE
FRANK BERRY
FRANK POSEY
FRANK HAWKE
FRANCES DEE
FRANK LEIGH
FRANCES DAY
FRANK FOSTER
FRANK HORNE
FRANK TOMEI
(10 rows)

Time: 2.960 ms

# Tags et tableaux



# PostgreSQL sait utiliser des tableaux

```
select tt.tid,  
       array_agg(tags.rowid) tags  
  from tags  
    join tid_tag tt  
      on tags.rowid = tt.tag  
 group by tt.tid  
 limit 3;
```

tid	tags
1	{1,2}
2	{3,4}
3	{5,6,7,8}

(3 rows)


time: 942.074 ms

# Types de données composites

<?xml?>

{JSON}

# SQL en 2014 : les *window functions*



# Finding the last counter value before reset

*Write some SQL here*

tick	nb	max
1	0	
2	10	
3	20	
4	30	
5	40	40
6	0	
7	20	
8	30	
9	60	60

(9 rows)

## Window Functions: lead() over()

```
select tick,
       nb,
       lead(nb) over (order by tick)
  from measures;
```

tick	nb	lead
1	0	10
2	10	20
3	20	30
4	30	40
5	40	0
6	0	20
7	20	30
8	30	60
9	60	

(9 rows)


# Window Functions and CASE

```
select tick, nb,
       case when lead(nb) over w < nb
             then nb
             when lead(nb) over w is null
             then nb
             else null
         end as max
  from measures
 window w as (order by tick);
```

tick	nb	max
1	0	
2	10	
3	20	
4	30	
5	40	40
6	0	
7	20	
8	30	
9	60	60

(9 rows)

# Un histogramme tout en SQL



# Histogramme tout en SQL

```
with drb_stats as (
    select min(drb) as min,
           max(drb) as max
      from team_stats
),
      histogram as (
        select width_bucket(drb, min, max, 9) as bucket,
               int4range(min(drb), max(drb), '[]') as range,
               count(*) as freq
      from team_stats, drb_stats
     group by bucket
    order by bucket
)
select bucket, range, freq,
       repeat('*', (freq::float / max(freq) over() * 30)::int)
  from histogram;
```

# Histogramme tout en SQL

bucket	range	freq	bar
1	[10,15)	52	
2	[15,20)	1363	**
3	[20,25)	8832	*****
4	[25,30)	20917	*****
5	[30,35)	20681	*****
6	[35,40)	9166	*****
7	[40,45)	2093	***
8	[45,50)	247	
9	[50,54)	20	
10	[54,55)	1	


(10 rows)

Time: 53.570 ms


# Savoir contourner les ORMs avec les wCTE

```
with queue as (
    insert into queue (extension)
        select id
            from extension
            where shortname = $1
    returning id, extension
)
select q.id, e.id as ext_id,
    e.fullname, e.uri, e.description
from      queue q
    join extension e on q.extension = e.id;
```



# Les Indexes dans PostgreSQL




# Les Indexes dans PostgreSQL




R-tree Hierarchy



# Les Indexes dans PostgreSQL



# Les jointures avec PostgreSQL




# Les jointures avec PostgreSQL

```
WITH upd AS (
    UPDATE target t
        SET counter = t.counter + s.counter,
        FROM source s
    WHERE t.id = s.id
    RETURNING s.id
)
INSERT INTO target(id, counter)
    SELECT id, sum(counter)
        FROM source s LEFT JOIN upd t USING(id)
    WHERE t.id IS NULL
    GROUP BY s.id
    RETURNING t.id
```



# POSTGRESQL IS YESQL!

Postgre**SQL**



# Questions?

Now is the time to ask!

