

# Writing & using Postgres Extensions

Postgre Open, Chicago, 2014

Dimitri Fontaine `dimitri@2ndQuadrant.fr`

Sept 17, 2014



## 2ndQuadrant France PostgreSQL Major Contributor

- pgloader, prefix, skytools, ...
- [apt.postgresql.org](http://apt.postgresql.org)
- CREATE EXTENSION
- CREATE EVENT TRIGGER
- MySQL migration tool, new pgloader version



# Writing & using Postgres Extensions

## Agenda

- How PostgreSQL extensibility works
- Things you can do with a PostgreSQL Extension
- The PostgreSQL indexing Framework
- **How to solve some practical use cases with existing extensions**
- Developping a new extension



# PostgreSQL Extensibility

PostgreSQL is *highly* **extensible**



# PostgreSQL Extensibility

```
select col1, col2 from table where col1 = 'something';
```



# PostgreSQL Extensibility

```
SELECT col  
  FROM table  
 WHERE stamped > date 'today' - interval '1 day';
```



# PostgreSQL Extensibility

```
select iprange, locid
  from geolite.blocks
 where iprange >= '91.121.37.122';
```

iprange	locid
91.121.0.0-91.121.159.255	75

(1 row)

Time: 1.220 ms



# PostgreSQL Extensibility: Operator Classes

SQL Operators are all dynamic and found in the catalogs

```
select amopopr::regoperator
  from pg_opclass c
 join pg_am am
    on am.oid = c.opcmethod
 join pg_amop amop
    on amop.amopfamily = c.opcfamily
 where opcintype = 'ip4r'::regtype
    and am.amname = 'gist';
```

```
amopopr
-----
>>=(ip4r,ip4r)
<<=(ip4r,ip4r)
>>(ip4r,ip4r)
<<(ip4r,ip4r)
&&(ip4r,ip4r)
=(ip4r,ip4r)
(6 rows)
```





# PostgreSQL Extensibility

```
select id, name, pos,  
       pos <@> point(-6.25,53.34) as miles  
from pubnames  
order by pos <-> point(-6.25,53.34)  
limit 10;
```



# PostgreSQL is Extensible

PostgreSQL plugins are data types and index support


- Data Type
- Input/Output functions
- Casts
- Operator Classes



# PostgreSQL is Extensible

PostgreSQL support several kind of indexes


- BTree, binary tree
- GiST, Generalized Search Tree
- SP-GiST, Space Partitioned GiST
- GIN, Generalized Inverted Index



# Binary Tree

Btree, the default index type

- Built for speed
- *unique* concurrency tricks
- Balanced
- support function: `cmp`
- operators: `<=` `<` `=` `>` `>=`




# Generalized Index Search Tree

## GiST or the Indexing API

- Built for comfort
- Balanced
- API: consistent, same, union
- API: penalty, picksplit
- API: compress, decompress
- operators: @> <@ && @@ = &< &>  
<<| ...


R-tree Hierarchy




# Generalized Inverted iNdex

Indexing several pointers per value, inversed cardinality

- Built for Text Search and Arrays
- Balanced
- API: compare, consistent
- API: extractValue, extractQuery
- operators: @> <@ && =



# Extensions and data types



# Some extensions example


46 Contribs, Community extensions, Private ones...

- **hll**
- **cube**
- **ltree**
- **citext**
- **hstore**
- **earthdistance**
- **pgq**
- **pg\_trgm**
- **wildspeed**
- **plproxy**
- **PostGIS**
- **ip4r**
- **intarray**
- **prefix**
- **pgfincore**
- **pgcrypto**
- **pgstattuple**
- **pg\_buffercache**
- **pg\_stat\_statements**
- **pgfincore**





# IP Ranges, ip4r



# IP Ranges, ip4r

```
table geolite.blocks limit 10;
```

iprange	locid
-----+-----	
1.0.0.0/24	17
1.0.1.0-1.0.3.255	49
1.0.4.0/23	14409
1.0.6.0/23	17
1.0.8.0/21	49
1.0.16.0/20	14614
1.0.32.0/19	47667
1.0.64.0/18	111
1.0.128.0-1.0.147.255	209
1.0.148.0/24	22537
(10 rows)	



# IP Ranges, ip4r, Geolocation

PostgreSQL allows using SQL and JOINS to match IP4R with geolocation.

```
select *  
  from geolite.blocks  
 join geolite.location  
   using(locid)  
 where iprange  
        >>=  
        '74.125.195.147';
```



# IP Ranges, ip4r, Geolocation

PostgreSQL allows using SQL and JOINS to match IP4R with geolocation.


```
select *
  from geolite.blocks
 join geolite.location
    using(locid)
 where iprange
        >=>
        '74.125.195.147';
```

-[ RECORD 1 ]-----	
locid	2703
iprange	74.125.189.24-74.125.
country	US
region	CA
city	Mountain View
postalcode	94043
location	(-122.0574,37.4192)
metrocode	807
areacode	650

Time: 1.335 ms




# Earth Distance



# How Far is The Nearest Pub

The point datatype is in-core

```
# CREATE TABLE pubnames
(  
  id      bigint,  
  pos     POINT,  
  name    text  
);
```



# How Far is The Nearest Pub

```
select name, pos
  from pubnames
order by pos <-> point(-6.25,53.346)
limit 3;
```

Pub Name	pos
Ned's	(-6.2519967,53.3458267)
Sub Lounge	(-6.2542332,53.3469085)
O'Neill's of Pearse Street	(-6.2524389,53.3448589)

(3 rows)

Time: 18.679 ms



# How Far is The Nearest Pub

```
CREATE INDEX on pubnames USING GIST(pos);
```

```
select name,  
       pos  
  from pubnames  
 order by pos <-> point(-0.12,51)  
  limit 3;
```

name	pos
Ned's	(-6.25,53.34)
Sub Lo	(-6.25,53.34)
O'Neil	(-6.25,53.34)

(3 rows)

Time: 0.849 ms





# How Far is The Nearest Pub, in Miles please.

```
create extension cube;  
create extension earthdistance;
```

```
select name,  
       pos <@> point(-6.25,53.34) miles  
       from pubnames  
order by pos <-> point(-6.25,53.34)  
limit 3;
```

name	miles
Ned's	0.06
Sub Lo	0.07
O'Neil	0.12

(3 rows)

Time: 1.335 ms



## Some pubs far away from here...

```
select c.name as city,  
pos <@> point(-6.25,53.34) as miles  
from pubnames p,  
    lateral (select name  
              from cities c  
              order by c.pos <-> p.pos  
              limit 1) c  
order by pos <-> point(-6.25,53.34)  
        desc  
limit 5;
```

city		miles
-----+-----		
Canterbury		399.44
Canterbury		378.91
Canterbury		392.08
Canterbury		397.30
Canterbury		379.68

(5 rows)

Time: 636.445 ms



# Geolocation: ip4r meets earthdistance



## Some pubs nearby... some place...

```
with geoloc as
(
  select location as l
    from location
   join blocks using(loid)
  where iprange
        >>=
        '212.58.251.195'
)
select name,
       pos <@> 1 miles
  from pubnames, geoloc
order by pos <-> 1
limit 10;
```

name	miles
Blue Anchor	0.299
Dukes Head	0.360
Blue Ball	0.337
Bell (aka The Rat)	0.481
on the Green	0.602
Fox & Hounds	0.549
Chequers	0.712
Sportsman	1.377
Kingswood Arms	1.205
Tattenham Corner	2.007

(10 rows)

Time: 3.275 ms



# Trigrams



# Trigrams and similarity

*similar* but not quite *like* the same

```
create extension pg_trgm;
```

```
select show_trgm('tomy') as tomy,  
       show_trgm('Tomy') as "Tomy",  
       show_trgm('tom torn') as "tom torn",  
       similarity('tomy', 'tom'),  
       similarity('dim', 'tom');
```

```
-[ RECORD 1 ]-----  
tomy          | {" t", " to", "my ",omy,tom}  
Tomy          | {" t", " to", "my ",omy,tom}  
tom torn      | {" t", " to", "om ",orn,"rn ",tom,tor}  
similarity    | 0.5  
similarity    | 0
```



# Trigrams and typos

Use your data to help your users out

```
select actor
  from products
 where actor ~* 'tomy';
actor
```

-----

(0 rows)

Time: <unregistered>

```
select actor
  from products
 where actor % 'tomy';
actor
```

-----

TOM TORN

TOM DAY


(2 rows)

Time: 26.972 ms



# Trigrams search indexing

```
create index on products using gist(actor gist_trgm_ops);
```



```
select actor  
  from products  
 where actor % 'tomy';  
      actor
```

```
-----  
TOM TORN  
TOM DAY  
(2 rows)
```

```
Time: 2.695 ms
```





# Trigrams and autocompletion

Use your data to help your users out

```
explain (costs off)
  select * from products where actor ~* 'tomy';
               QUERY PLAN
```

---

```
Index Scan using products_actor_idx on products
  Index Cond: ((actor)::text ~* 'tomy'::text)
(2 rows)
```



# Trigrams and autocompletion

Use your data to help your users out

```
select actor
  from products
 where actor % 'fran'
order by actor <-> 'fran'
 limit 10;
```

actor

---

```
FRANK HAWKE
FRANK BERRY
FRANK POSEY
FRANK HAWKE
FRANCES DEE
FRANK LEIGH
FRANCES DAY
FRANK FOSTER
FRANK HORNE
FRANK TOMEI
(10 rows)
```

Time: 2.960 ms



# Advanced Array Indexing with intarray



## Last.fm allows users to tag tracks

	tag	n
	the brian setzer orchestra	1
select t.tag,	setzer	13
count(tt.tid) n	rockabilly setzer style	4
from tid_tag tt	setzer is a true guitarhero	9
join tags t	brian setzer orchestra	3
on tt.tag = t.rowid	brian setzer is god	1
where t.tag ~* 'setzer'	brian setzer	1
group by t.tag;	brain setzer orchestra	2
	(8 rows)	

time: 644.826 ms



# Last.fm allows users to tag tracks

```
create extension intarray;
```

```
select tt.tid,  
       array_agg(tags.rowid) tags  
from tags  
join tid_tag tt  
  on tags.rowid = tt.tag  
group by tt.tid  
limit 3;
```

tid	tags
1	{1,2}
2	{3,4}
3	{5,6,7,8}

(3 rows)

time: 942.074 ms



## Prepare for intarray indexing

Denormalize the data set thanks to PostgreSQL Arrays

```
create table track_tags as
  select tt.tid, array_agg(tags.rowid) as tags
  from tags join tid_tag tt on tags.rowid = tt.tag
  group by tt.tid;

create index on track_tags using gin(tags gin__int_ops);
```



# Search for several tags at once

## Intersection of multiple criteria

```
select array_agg(rowid)
  from tags
 where tag = 'blues'
        or tag = 'rhythm and blues';
```

array\_agg

-----  
{3,739}

(1 row)

time: 0.684 ms



# The query\_int data type

intarray has powerful indexing and searching facilities

```
select format('(%s)',  
            array_to_string(  
                array_agg(rowid), '&')  
           )::query_int as query  
from tags  
where tag = 'blues'  
       or tag = 'rhythm and blues';
```

```
query  
-----  
3 & 739  
(1 row)
```

time: 0.747 ms





## Putting it all together

```
with t(query) as (  
  select format('(%s)',  
    array_to_string(  
      array_agg(rowid), '&')  
    )::query_int as query  
  from tags  
  where tag = 'blues'  
    or tag = 'rhythm and blues'  
)  
select track.tid  
  from track_tags tt  
  join tids track  
    on tt.tid = track.rowid, t  
  where tt.tags @@ t.query  
 limit 10;
```

tid
TRCJLCC12903CBF4AE
TRCIF0V128F92F6F4C
TRCYUVJ128F425C8F1
TRCNTF0128F92F6564
TRCDRGT12903CE64BF
TRCWAED128F42A837B
TRCWFEM128F9320F94
TRCQCQH128F932E707
TRCUMTA12903CD67EE
TRJJYUT12903CFB13B

(10 rows)

Time: 7.630 ms





# HStore: NoSQL meets PostgreSQL

- Key-Value Store
- Denormalisation
- Indexing (GiST, GIN)
- Operators
- SQL



# HStore: NoSQL meets PostgreSQL

```
create table preferences
(  
  email          text          primary key,  
  language       text,  
  timezone       text,  
  properies      hstore  
);
```



# HStore: NoSQL meets PostgreSQL

```
INSERT INTO preferences
```

```
VALUES
```

```
('dimitri@2ndQuadrant.fr', 'fr_FR', 'Europe/Paris',  
 'skills => PostgreSQL,  
  Extensions => "prefix, base64, pgextwlist, preprepare",  
  Software => "pgloader");
```

```
('simon@2ndQuadrant.com', 'en_UK', 'Europe/London',  
 'skills => "PostgreSQL, Replication",  
  Software => "PostgreSQL, repmgr, pg_standby");
```

```
INSERT 0 2
```



# HStore: NoSQL meets PostgreSQL

```
~# select * from preferences;
-[ RECORD 1 ]-----
email      | dimitri@2ndQuadrant.fr
language   | fr_FR
timezone   | Europe/Paris
properties | "skills"=>"PostgreSQL",
                  "Software"=>"pgloader",
                  "Extensions"=>"prefix, base64, pgextwlist, preprep
-[ RECORD 2 ]-----
email      | simon@2ndQuadrant.com
language   | en_UK
timezone   | Europe/London
properties | "skills"=>"PostgreSQL, Replication",
                  "Software"=>"PostgreSQL, repmgr, pg_standby"
```



# HStore: NoSQL meets PostgreSQL

```
~# select email  
      from preferences  
     where properties ? 'Extensions';  
           email
```

```
-----  
dimitri@2ndQuadrant.fr  
(1 row)
```



# HStore: NoSQL meets PostgreSQL

```
~# select email  
    from preferences  
   where (properties -> 'skills') ~ 'PostgreSQL';  
        email
```

```
-----  
dimitri@2ndQuadrant.fr  
simon@2ndQuadrant.com  
(2 rows)
```





# HStore and Parametrized Triggers



## Compute *duration* in a before trigger

We need a table and some data

```
create table foo
(
  id          serial          primary key,
  d_start     timestamptz default now(),
  d_end       timestamptz,
  duration    interval
);

insert into foo(d_start, d_end)
select now() - 10 * random() * interval '1 min',
       now() + 10 * random() * interval '1 min'
from generate_series(1, 10);
```



# Populating an hstore from a record

Filling in a column when the name is a parameter

```
select (foo #>= hstore('duration', '10 mins')).*  
  from foo  
 limit 1;
```

```
-[ RECORD 1 ]-----  
id           | 1  
d_start      | 2013-10-24 18:57:36.725212+02  
d_end        | 2013-10-24 19:04:56.171473+02  
duration     | 00:10:00
```



## The hstore based *parametrized* trigger

```
create or replace function tg_duration()
  returns trigger
  language plpgsql
as '
declare
  hash hstore := hstore(NEW);
  duration interval;
begin
  duration := (hash -> TG_ARGV[1])::timestampz
              - (hash -> TG_ARGV[0])::timestampz;

  NEW := NEW #= hstore(TG_ARGV[2], duration::text);

  RETURN NEW;
end;
';
```



# Installing the trigger

To be able to modify what's inserted we need a before trigger

```
create trigger compute_duration
  before insert on foo
  for each row
execute procedure tg_duration('d_start',
                              'd_end',
                              'duration');
```



## Using the trigger: watch the duration

```
truncate foo;
insert into foo(d_start, d_end)
  select now()
    - 10 * random()
      * interval '1 min',
    now()
    + 10 * random()
      * interval '1 min'
  from generate_series(1, 10);
```


```
select duration
  from foo;
      duration
```

```
-----
00:03:48.003135
00:10:57.727407
00:01:13.637183
00:10:33.820578
00:13:11.607287
00:04:41.224213
00:08:26.842229
00:12:16.630843
00:09:51.418547
00:08:52.968195
(10 rows)
```





PL/Proxy is all about *Sharding*



We're going to use it for *Remote Procedure Call*





# Classic Auditing

```
create table example
(
    id serial,
    f1 text,
    f2 text
);
```

```
create table audit
(
    change_date timestampz
                default now(),
    before hstore,
    after  hstore
);
```



# Classic trigger based Auditing

```
~# begin;
~*# update example set f1 = 'b' where id = 1;
~*# rollback;

~# select * from audit;
  change_date | before | after
-----+-----+-----
(0 rows)
```



# Setting up PL/Proxy

```
create extension plproxy;

create server local
foreign data wrapper plproxy
options (p0 'dbname=dim');

create user mapping
for public
server local
options (user 'dim');
```



## PL/Proxy: Basic Testing

```
create function test_proxy
    (i int)
    returns int
    language plproxy
as '
    cluster ''local'';
    select i;
';
```

```
select test_proxy(1);
test_proxy
-----
          1
(1 row)

Time: 0.866 ms
```



# Implementing Autonomous Transactions for Auditing



## Trigger Functions 1/3: the trigger

```
create function audit_trigger()  
  returns trigger  
  language plpgsql  
as '  
begin  
  perform audit_proxy(old, new);  
  return new;  
end;  
';
```



## Trigger Functions 2/3: the proxy

```
create function audit_proxy
(
    old example,
    new example
)
returns void
language plproxy
as '
    cluster ''local'';
    target audit;
';
```



## Trigger Functions 3/3: the implementation

```
create function audit
(
  old example,
  new example
)
returns void
language SQL
as '
  INSERT INTO audit(before, after)
    SELECT hstore(old), hstore(new);
';
```





# Trigger Definition

```
drop trigger if exists audit on example;

create trigger audit
  after update on example
    for each row
      -- careful, defaults to FOR EACH STATEMENT!
execute procedure audit_trigger();
```



# Autonomous Auditing Transaction

```
~# begin;
```

```
BEGIN
```

```
~*# update example set f1 = 'b' where id = 1;
```

```
UPDATE 1
```

```
~*# rollback;
```

```
ROLLBACK
```



# Autonomous Auditing Transaction

We did ROLLBACK; the transaction

```
~# select change_date,  
         before,  
         after,  
         after - before as diff  
       from audit;  
-[ RECORD 1 ]-----  
change_date | 2013-10-14 14:29:09.685105+02  
before      | "f1"=>"a", "f2"=>"a", "id"=>"1"  
after       | "f1"=>"b", "f2"=>"a", "id"=>"1"  
diff        | "f1"=>"b"
```



## State of The Art Cardinality Estimation Algorithm



# Creating the unique visitors tracking table

```
CREATE EXTENSION hll;

-- Create the destination table
CREATE TABLE daily_uniques (
    DATE            DATE UNIQUE,
    users           hll
);

-- Our first aggregate update
UPDATE daily_uniques
    SET users = hll_add(users,
                        hll_hash_text('123.123.123.123'))
WHERE date = current_date;
```



# Production ready updates

```
--  
-- First upload a new batch, e.g. using  
-- CREATE TEMP TABLE new_batch as VALUES(), (), ...;  
--  
WITH hll(agg) AS (  
    SELECT hll_add_agg(hll_hash_text(value))  
        FROM new_batch  
)  
UPDATE daily_uniques  
    SET users = CASE WHEN hll.agg IS NULL THEN users  
                    ELSE hll_union(users, hll.agg)  
                END  
    FROM hll  
WHERE date = current_date;
```



# Daily Reporting

```
with stats as (  
    select date,  
           #users as daily,  
  
           #hll_union_agg(users)  
           over() as total  
    from daily_uniques  
)  
select date,  
       daily,  
       daily/total*100  
    from stats  
order by date;
```

date	daily	percent
2013-02-22	401677	25.19
2013-02-23	660187	41.41
2013-02-24	869980	54.56
2013-02-25	154996	9.72

(4 rows)



# Monthly Reporting


```
select to_char(date, 'YYYY/MM'),  
       #hll_union_agg(users)  
from daily_uniques  
group by 1;
```

month	monthly
2013/02	1960380

(1 row)




## New in 9.3: Background Workers



# New in 9.3: Background Workers

Start autonomous user processes within the database server

- Job Scheduler (autovacuum like maintenance)
- PGQ Ticker
- Replication Tasks
- Parallel Queries



# Background Workers C API

```
void _PG_init(void)
{
    BackgroundWorker worker;

    worker.bgw_flags = BGWORKER_SHMEM_ACCESS |
        BGWORKER_BACKEND_DATABASE_CONNECTION;
    worker.bgw_start_time    = BgWorkerStart_RecoveryFinished;
    worker.bgw_main          = worker_spi_main;
    worker.bgw_sighup        = worker_spi_sighup;
    worker.bgw_sigterm       = worker_spi_sigterm;
    worker.bgw_name          = "count relations";
    worker.bgw_restart_time  = BGW_NEVER_RESTART;
    worker.bgw_main_arg      = NULL;
    RegisterBackgroundWorker(&worker);

    BackgroundWorkerInitializeConnection("dbname", "username");
}
```

# Background Workers C API

`bgw_start_time`

- `BgWorkerStart_PostmasterStart`
- `BgWorkerStart_ConsistentState`
- `BgWorkerStart_RecoveryFinished`



# Background Workers and SPI

```
StartTransactionCommand();  
SPI_connect();  
PushActiveSnapshot(GetTransactionSnapshot());  
  
/* build query, might be static string */  
  
SPI_execute(query, true/false /* read only */, 0);  
  
/* process query results */  
  
SPI_finish();  
PopActiveSnapshot();  
CommitTransactionCommand();
```



## New in 9.3: Background Workers

Can request *shared memory*, server must be restarted.

```
shared_preload_libraries = 'count_relations'
```




Extensions: Let's make a new one!



# A new integer data type: base36

Internally store `bigint`, reuse *internals*



Only code the *input/output* functions, in C





# A new integer data type: base36

```
create extension base36;
```

i	x
<hr/>	
0	0
1	1
2	2
3	3
4	4
5	5
6	6
7	7
8	8
9	9
10	A

i	x
<hr/>	
10000	7PS
10001	7PT
10002	7PU
10003	7PV
10004	7PW
10005	7PX
10006	7PY
10007	7PZ
10008	7Q0
10009	7Q1
10010	7Q2

i	x
<hr/>	
100000000	1NJCHS
100000001	1NJCHT
100000002	1NJCHU
100000003	1NJCHV
100000004	1NJCHW
100000005	1NJCHX
100000006	1NJCHY
100000007	1NJCHZ
100000008	1NJCIO
100000009	1NJCI1
100000010	1NJCI2

# Input Output Functions

```
CREATE OR REPLACE FUNCTION base36_in(cstring)
RETURNS base36
AS '$libdir/base36'
LANGUAGE C IMMUTABLE STRICT;
```

```
CREATE OR REPLACE FUNCTION base36_out(base36)
RETURNS cstring
AS '$libdir/base36'
LANGUAGE C IMMUTABLE STRICT;
```



# Input Output Functions

```
CREATE OR REPLACE FUNCTION base36_recv(internal)
RETURNS base36
AS '$libdir/base36'
LANGUAGE C IMMUTABLE STRICT;
```

```
CREATE OR REPLACE FUNCTION base36_send(base36)
RETURNS bytea
AS '$libdir/base36'
LANGUAGE C IMMUTABLE STRICT;
```



## Some C code

```
#include "postgres.h"

#ifdef PG_VERSION_NUM
#error "Unsupported too old PostgreSQL version"
#endif

#if PG_VERSION_NUM / 100 != 903 \
    && PG_VERSION_NUM / 100 != 904
#error "Unknown or unsupported PostgreSQL version"
#endif

PG_MODULE_MAGIC;
```



## Some C code

```
static inline
base36 base36_from_str(const char *str)
{
    /* ... C code here ... */
}
```

```
static inline
char *base36_to_str(base36 c)
{
    /* ... C code here ... */
}
```



# Interfacing with PostgreSQL

```
Datum base36_in(PG_FUNCTION_ARGS);  
Datum base36_out(PG_FUNCTION_ARGS);  
Datum base36_recv(PG_FUNCTION_ARGS);  
Datum base36_send(PG_FUNCTION_ARGS);  
Datum base36_cast_to_text(PG_FUNCTION_ARGS);  
Datum base36_cast_from_text(PG_FUNCTION_ARGS);  
Datum base36_cast_to_bigint(PG_FUNCTION_ARGS);  
Datum base36_cast_from_bigint(PG_FUNCTION_ARGS);
```



# Interfacing with PostgreSQL

```
PG_FUNCTION_INFO_V1(base36_in);
Datum
base36_in(PG_FUNCTION_ARGS)
{
    char *str = PG_GETARG_CSTRING(0);
    PG_RETURN_INT64(base36_from_str(str));
}
```

```
PG_FUNCTION_INFO_V1(base36_out);
Datum
base36_out(PG_FUNCTION_ARGS)
{
    base36 c = PG_GETARG_INT64(0);
    PG_RETURN_CSTRING(base36_to_str(c));
}
```



# CREATE TYPE

```
CREATE TYPE base36 (  
    INPUT          = base36_in,  
    OUTPUT         = base36_out,  
    RECEIVE        = base36_recv,  
    SEND           = base36_send,  
    LIKE           = bigint,  
    CATEGORY       = 'N'  
);  
COMMENT ON TYPE base36  
    IS 'bigint written in base36: [0-9A-Z]+';
```





# A minimum amount of CAST

```
CREATE FUNCTION text(base36)
RETURNS text
AS '$libdir/base36',
    'base36_cast_to_text'
LANGUAGE C IMMUTABLE STRICT;

CREATE CAST (text as base36)
WITH FUNCTION base36(text)
AS IMPLICIT;
```

```
CREATE CAST (base36 as text)
WITH FUNCTION text(base36);

CREATE CAST (bigint as base36)
WITHOUT FUNCTION
AS IMPLICIT;

CREATE CAST (base36 as bigint)
WITHOUT FUNCTION
AS IMPLICIT;
```



## Reuse *internals*: comparison functions

```
CREATE OR REPLACE FUNCTION base36_eq(base36, base36)
RETURNS boolean LANGUAGE internal IMMUTABLE AS 'int8eq';
```

```
CREATE OR REPLACE FUNCTION base36_ne(base36, base36)
RETURNS boolean LANGUAGE internal IMMUTABLE AS 'int8ne';
```

```
CREATE OR REPLACE FUNCTION base36_lt(base36, base36)
RETURNS boolean LANGUAGE internal IMMUTABLE AS 'int8lt';
```

```
CREATE OR REPLACE FUNCTION base36_le(base36, base36)
RETURNS boolean LANGUAGE internal IMMUTABLE AS 'int8le';
```



## Reuse *internals*: comparison functions

```
CREATE OR REPLACE FUNCTION base36_gt(base36, base36)
RETURNS boolean LANGUAGE internal IMMUTABLE AS 'int8gt';
```

```
CREATE OR REPLACE FUNCTION base36_ge(base36, base36)
RETURNS boolean LANGUAGE internal IMMUTABLE AS 'int8ge';
```

```
CREATE OR REPLACE FUNCTION base36_cmp(base36, base36)
RETURNS integer LANGUAGE internal IMMUTABLE AS 'btint8cmp';
```



# Register operators

```
CREATE OPERATOR = (  
    LEFTARG = base36,  
    RIGHTARG = base36,  
    PROCEDURE = base36_eq,  
    COMMUTATOR = '=',  
    NEGATOR = '<>',  
    RESTRICT = eqsel,  
    JOIN = eqjoinsel  
);  
COMMENT ON OPERATOR =(base36, base36) IS 'equals?';
```



## Add in btree index support, stolen from bigint

```
CREATE OPERATOR CLASS btree_base36_ops
DEFAULT FOR TYPE base36 USING btree
AS
```

```
OPERATOR          1          < ,
OPERATOR          2          <= ,
OPERATOR          3          = ,
OPERATOR          4          >= ,
OPERATOR          5          > ,
FUNCTION          1          base36_cmp(base36, base36);
```



# Packaging an extension



# Packaging an extension

We need a Makefile

```
EXTENSION    = base36
MODULES      = base36
DATA         = base36--1.0.sql base36.control
```

```
LDFLAGS=-lrt
```

```
PG_CONFIG ?= pg_config
PGXS = $(shell $(PG_CONFIG) --pgxs)
include $$(PGXS)
```



# Installing an extension

```
$ make install
```

```
/usr/local/bin/ccache /usr/bin/gcc -O2 -Wall -Wmissing-prototy  
/usr/local/bin/ccache /usr/bin/gcc -O2 -Wall -Wmissing-prototy  
/bin/sh /Users/dim/pgsql/ddl/lib/pgxs/src/makefiles/../../conf  
/bin/sh /Users/dim/pgsql/ddl/lib/pgxs/src/makefiles/../../conf  
/bin/sh /Users/dim/pgsql/ddl/lib/pgxs/src/makefiles/../../conf  
/usr/bin/install -c -m 644 base36.control '/Users/dim/pgsql/do  
/usr/bin/install -c -m 644 base36--1.0.sql base36.control '/Us  
/usr/bin/install -c -m 755 base36.so '/Users/dim/pgsql/ddl/li
```





## Enjoying our new extension

```
create extension base36;

create table demo(i bigint, x base36);

insert into demo(i, x)
  select n, n::bigint
  from generate_series(0, 10) t(n);

insert into demo(i, x)
  select n, n::bigint
  from generate_series(10000, 10010) t(n);

insert into demo(i, x)
  select n, n::bigint
  from generate_series(100000000, 100000010) t(n);

create index on demo(x);
```



# A new integer data type: base36

```
create extension base36;
```

i		x
<hr/>		
0		0
1		1
2		2
3		3
4		4
5		5
6		6
7		7
8		8
9		9
10		A

i		x
<hr/>		
10000		7PS
10001		7PT
10002		7PU
10003		7PV
10004		7PW
10005		7PX
10006		7PY
10007		7PZ
10008		7Q0
10009		7Q1
10010		7Q2

i		x
<hr/>		
1000000000		1NJCHS
1000000001		1NJCHT
1000000002		1NJCHU
1000000003		1NJCHV
1000000004		1NJCHW
1000000005		1NJCHX
1000000006		1NJCHY
1000000007		1NJCHZ
1000000008		1NJCIO
1000000009		1NJCI1
1000000010		1NJCI2



# PostgreSQL is YeSQL!



# Recap

We saw a number of extensions, each with a practical use case

**ip4r** IP Ranges and Geolocation

**Earth** Longitude, Latitude, Computing distances on a map

**Trigrams** Fixing typos, autocompletion

**Intarray** Indexing Tag Searches

**HStore** Schemaless development, Generic Auditing triggers

**PL/Proxy** Sharding, RPC, Autonomous Transactions

**HLL** Cardinalities, Unique Visitors

**BGWorkers** PostgreSQL managed processes

**base36** bigints with letters



# Questions?

Now is the time to ask!

